


Lecture 7

- Administration

Chapter 5

Monitors & Condition Synchronization



monitors & condition synchronization

Concepts: monitors:

encapsulated data + access procedures
mutual exclusion + condition synchronization
single access procedure active in the monitor
nested monitors

Models: guarded actions

Practice: private data and synchronized methods (exclusion). wait(), notify() and notifyAll() for condition synch. single thread active in the monitor at a time

Example of Notify()

P1 locks and puts and unlocks

P2 locks, tries to put, waits&unlock

P3 tries to put, waits&unlock

C1 locks

C2 blocked

C3 blocked

C1 calls notify(), unlocks

P2 is awakened, tries to put

C2 locks, tries to get, waits&unlock

C3 locks, tries to get, waits&unlock

P2 locks, and puts, calls notify(), unlocks

P3 is awakened, gets lock, waits again

There is no one to call notify() to wake up P3 or C2 or C3

Java Monitors

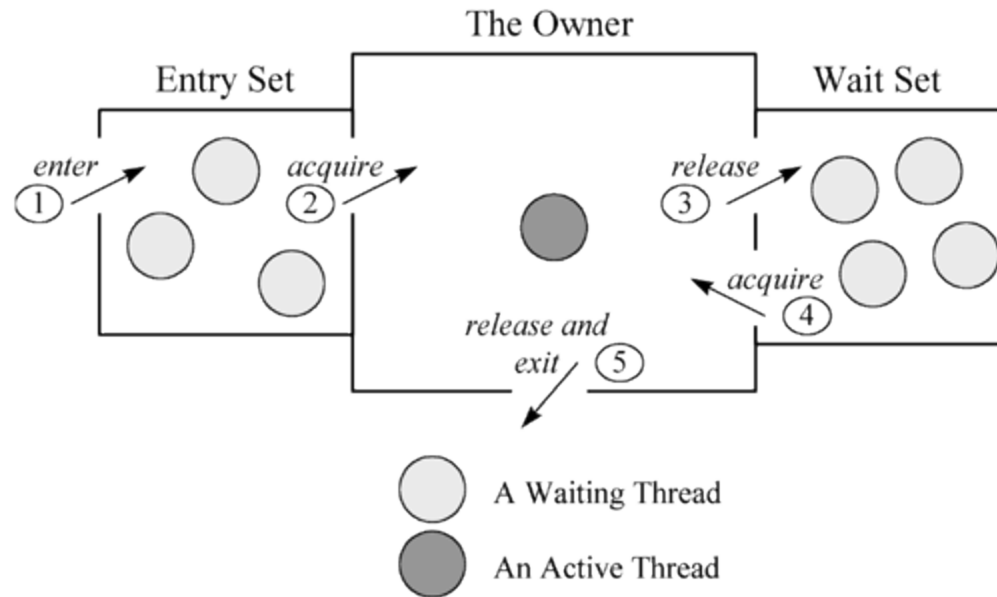
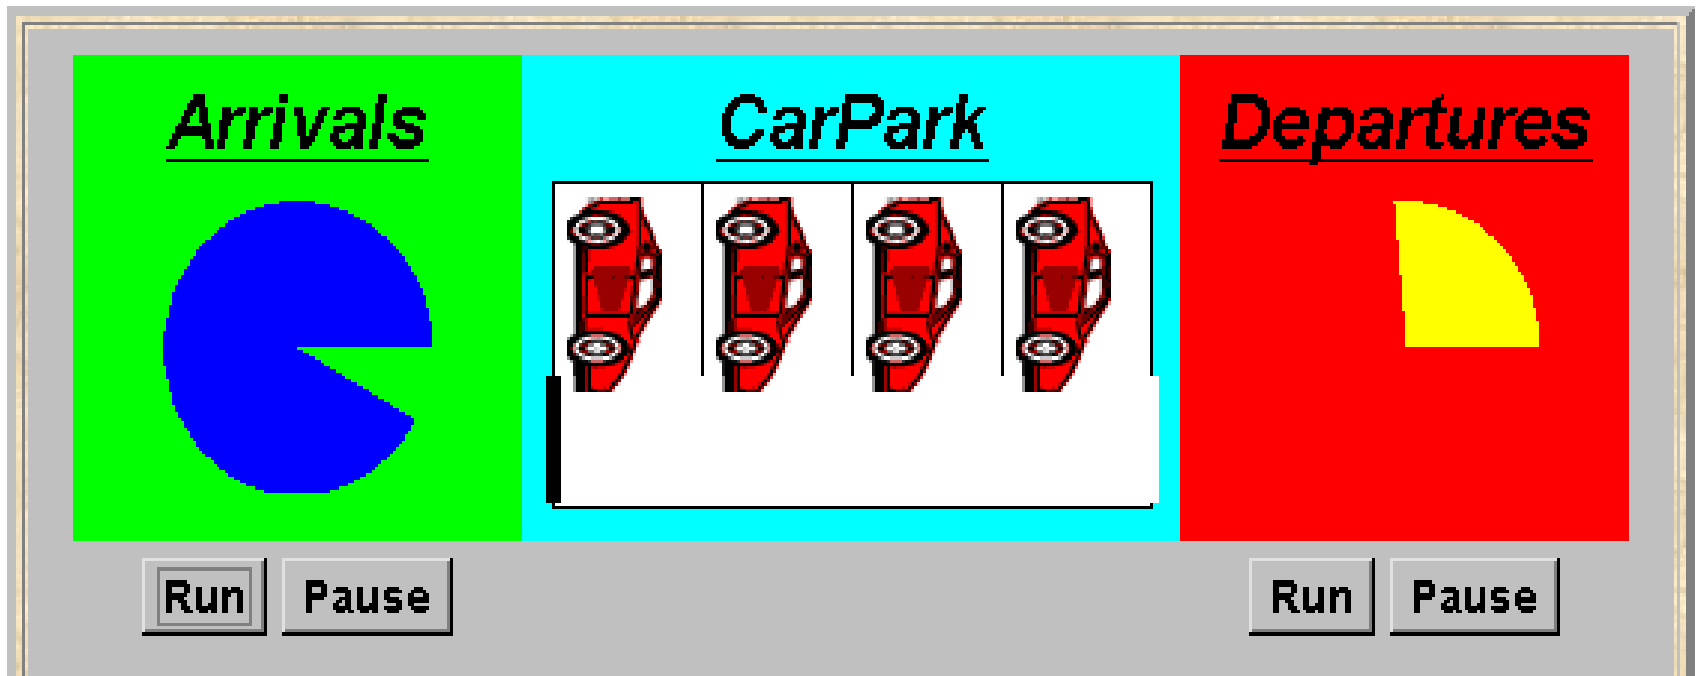


Figure 20-1. A Java monitor.

5.1 Condition synchronization



A controller is required for a carpark, which only permits cars to enter when the carpark is not full and does not permit cars to leave when there are no cars in the carpark. Car arrival and departure are simulated by separate threads.

carpark model


- ◆ Events or actions of interest?

arrive and depart

- ◆ Identify processes.

arrivals, departures and carpark control

- ◆ Define each process and interactions (structure).



carpark model

```
CARPARKCONTROL(N=4) = SPACES[N],  
SPACES[i:0..N] = (when(i>0) arrive->SPACES[i-1]  
                  | when(i<N) depart->SPACES[i+1]  
                  ).
```

```
ARRIVALS = (arrive->ARRIVALS).
```

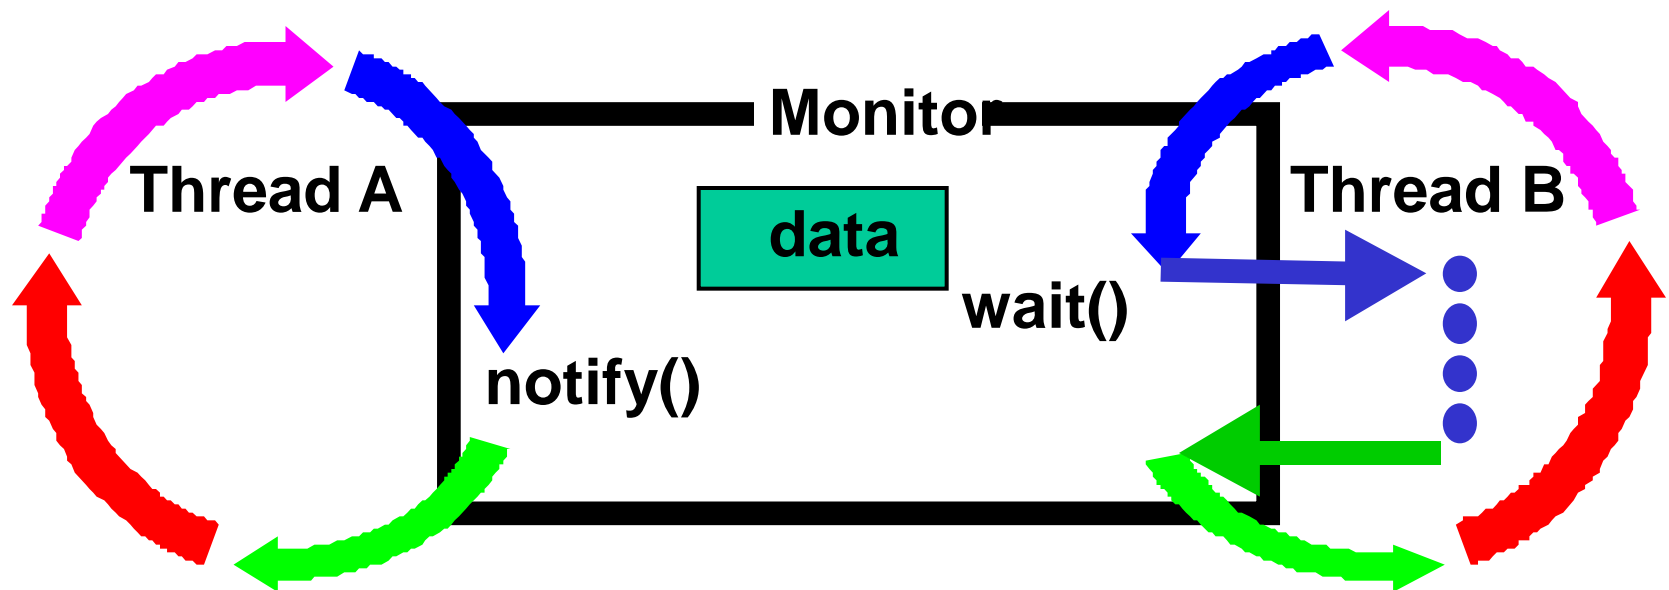
```
DEPARTURES = (depart->DEPARTURES).
```

```
|| CARPARK =  
    (ARRIVALS | | CARPARKCONTROL(4) | | DEPARTURES).
```


condition synchronization in Java

We refer to a thread ***entering*** a monitor when it acquires the mutual exclusion lock associated with the monitor and ***exiting*** the monitor when it releases the lock.

Wait() - causes the thread to exit the monitor, permitting other threads to enter the monitor.



Java Monitors

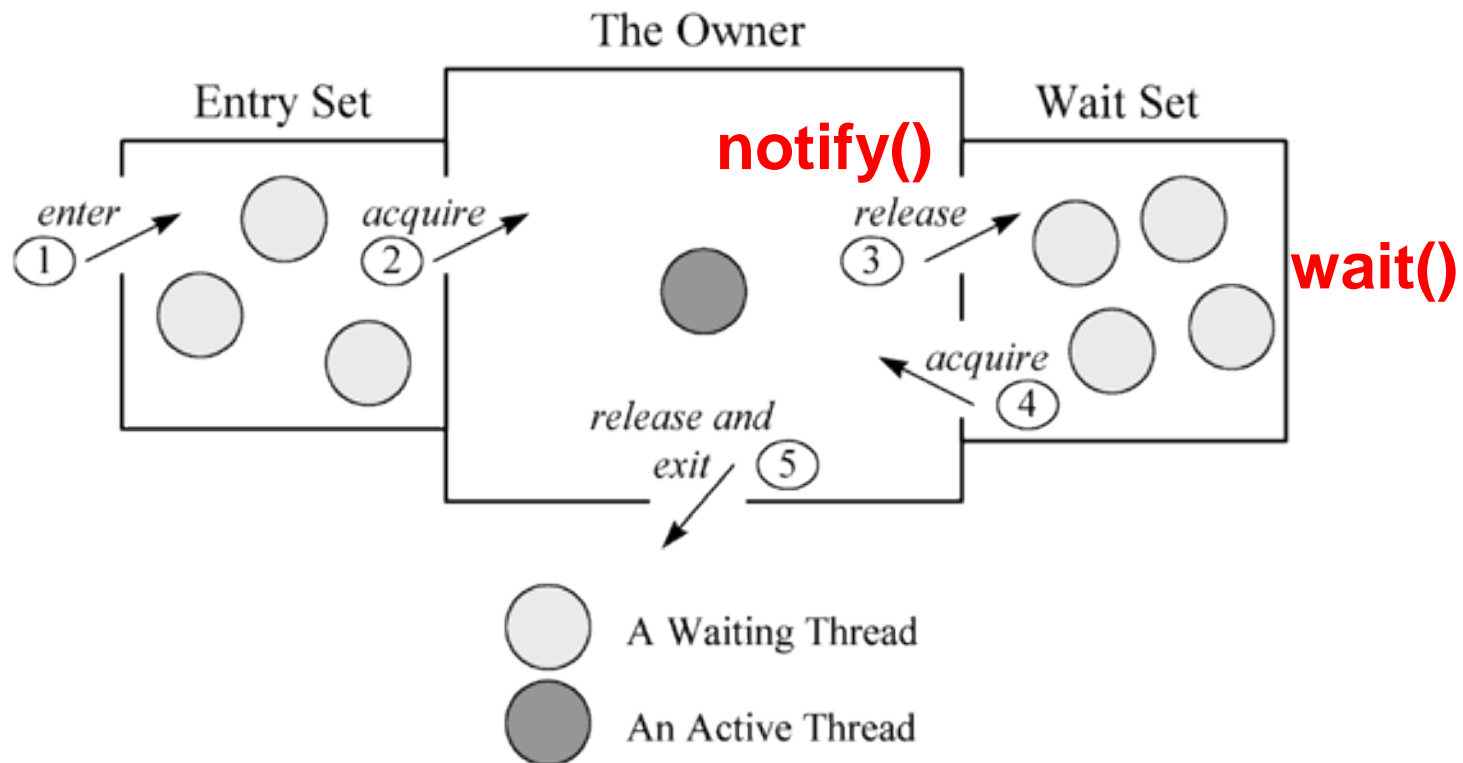


Figure 20-1. A Java monitor.

condition synchronization in Java

FSP: when *cond* act -> NEWSTAT

```
Java:   public synchronized void act()  
              throws InterruptedException  
      {  
          while (!cond) wait();  
          // modify monitor data  
          notifyAll()  
      }
```

The **while** loop is necessary to retest the condition *cond* to ensure that *cond* is indeed satisfied when it re-enters the monitor.

notifyall() is necessary to awaken other thread(s) that may be waiting to enter the monitor now that the monitor data has been changed.

Semaphores

Semaphores are widely used for dealing with inter-process synchronization in operating systems. Semaphore s is an integer variable that can take only non-negative values.

The only operations permitted on s are $up(s)$ and $down(s)$. Blocked processes are held in a FIFO queue.

```
down(s): if s > 0 then
           decrement s
        else
           block execution of the
           calling process

up(s):    if processes blocked on s
           then
           awaken one of them
        else
           increment s
```

modeling semaphores


To ensure analyzability, we only model semaphores that take a finite range of values. If this range is exceeded then we regard this as an **ERROR**. **N** is the initial value.

```
const Max = 3
range Int = 0..Max

SEMAPHORE(N=0) = SEMA[N],
SEMA[v:Int]    = (up->SEMA[v+1]
                  | when(v>0) down->SEMA[v-1]
                  ),
SEMA[Max+1]    = ERROR.
```

LTS?

modeling semaphores



Action **down** is only accepted when value v of the semaphore is greater than 0.

Action **up** is not guarded.

Trace to a violation:

$\text{up} \rightarrow \text{up} \rightarrow \text{up} \rightarrow \text{up}$

semaphore demo - model

Three processes $p[1..3]$ use a shared semaphore `mutex` to ensure mutually exclusive access (action critical) to some resource.

```
LOOP = (  
    mutex.down->critical->mutex.up->LOOP  
).  
  
|| SEMADEMO = (p[1..3]:LOOP  
    || {p[1..3]}::mutex:SEMAPHORE(1)).
```

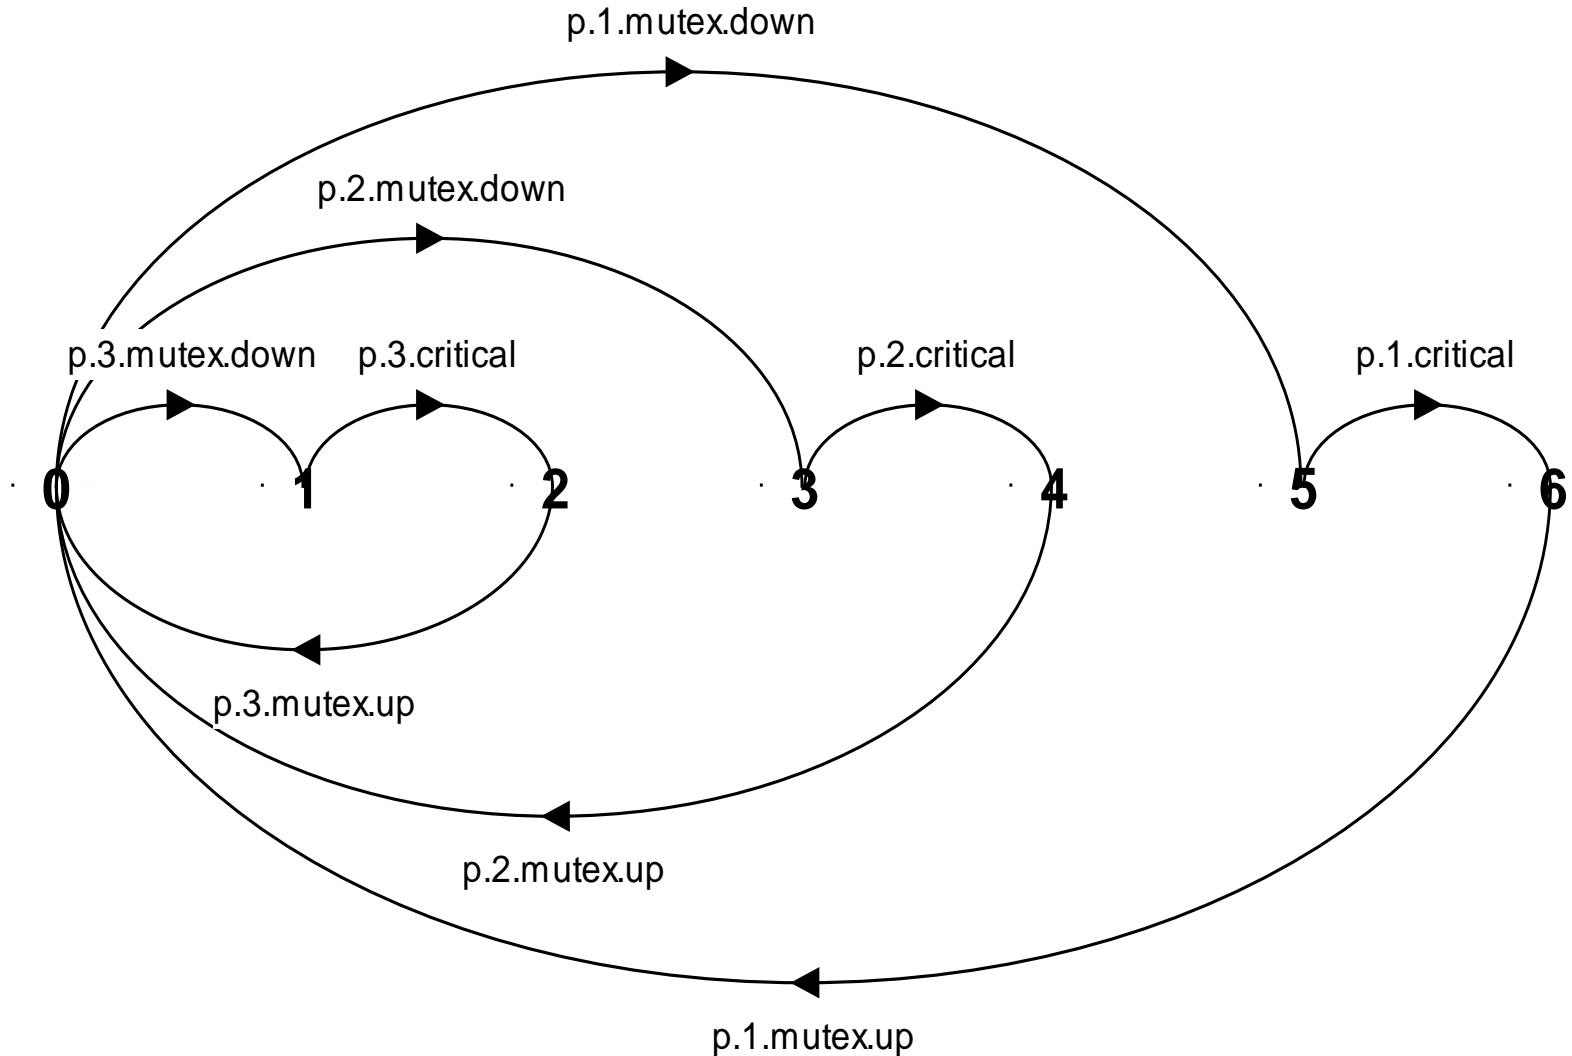
For mutual exclusion, the semaphore initial value is 1. *Why?*

*Is the **ERROR** state reachable for SEMADEMO?*

*Is a **binary** semaphore sufficient (i.e. $\text{Max}=1$)?*

LTS?

semaphore demo - model



semaphores in Java

Semaphores are passive objects, therefore implemented as **monitors**.

(In practice, semaphores are a low-level mechanism often used in implementing the higher-level monitor construct.)

```
public class Semaphore {
    private int value;

    public Semaphore (int initial)
        {value = initial;}

    synchronized public void up() {
        ++value;
        notifyAll();
    }

    synchronized public void down()
        throws InterruptedException {
        while (value== 0) wait();
        --value;
    }
}
```

Is it safe to use notify() here rather than notifyAll()?

SEMADEMO display

Mutex
0


current
semaphore
value

thread 1 is
executing
critical
actions.


thread 2 is
blocked
waiting.

thread 3 is
executing non-
critical
actions.


Thread 1



Thread 2



Thread 3



Run

Pause

Run

Pause

Run

Pause

SEMADEMO

*What if we adjust the time that each thread spends in its **critical section** ?*

◆ large resource requirement - *more conflict?*


(eg. more than 67% of a rotation)?

◆ small resource requirement - *no conflict?*

(eg. less than 33% of a rotation)?

Hence the time a thread spends in its critical section should be kept as short as possible.


5.3 Bounded Buffer



A bounded buffer consists of a fixed number of slots. Items are put into the buffer by a *producer* process and removed by a *consumer* process. It can be used to smooth out transfer rates between the *producer* and *consumer*.


(see car park example)

bounded buffer



The behaviour of BOUNDEDBUFFER is independent of the actual data values, and so can be modelled in a data-independent manner.

LTS:



Data-independent model

bounded buffer

```
BUFFER(N=5) = COUNT[0],  
COUNT[i:0..N]  
    = (when (i<N) put->COUNT[i+1]  
      |when (i>0) get->COUNT[i-1]  
      ).
```

```
PRODUCER = (put->PRODUCER).
```

```
CONSUMER = (get->CONSUMER).
```

```
||BOUNDEDBUFFER = (PRODUCER || BUFFER(5) || CONSUMER).
```