

# Lecture 9

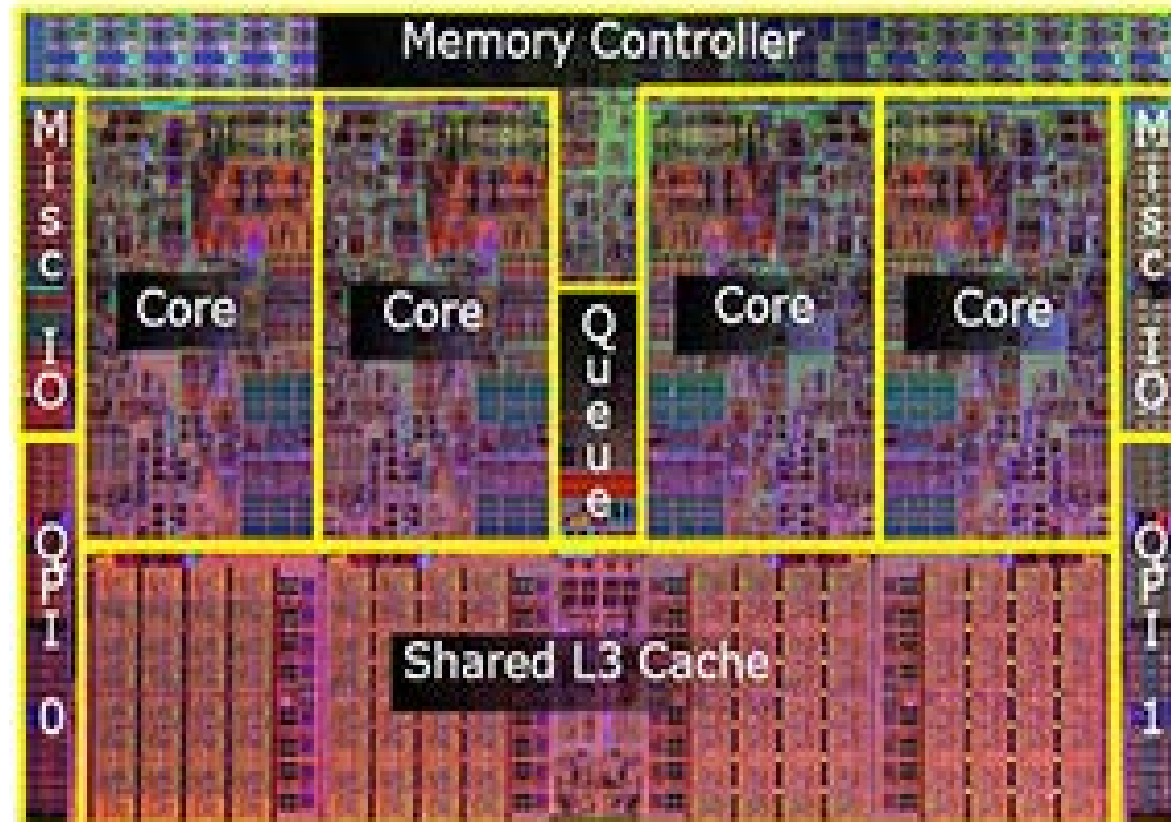
- ❑ Administration
- ❑ Pthreads
- ❑ Exceptions in C

<https://computing.llnl.gov/tutorials/pthreads/>

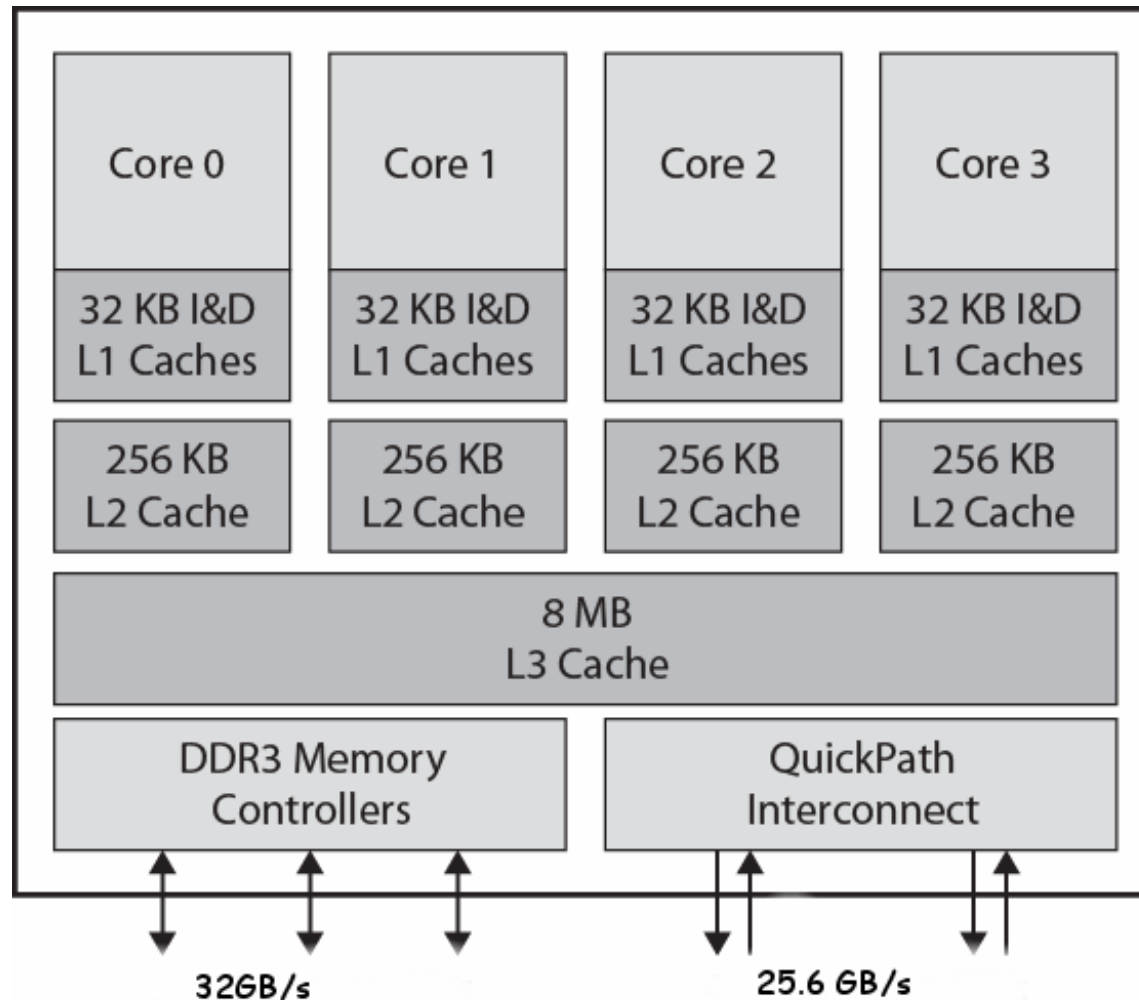
# Intel Core i7

approx 45x45 mm

45 nm feature size

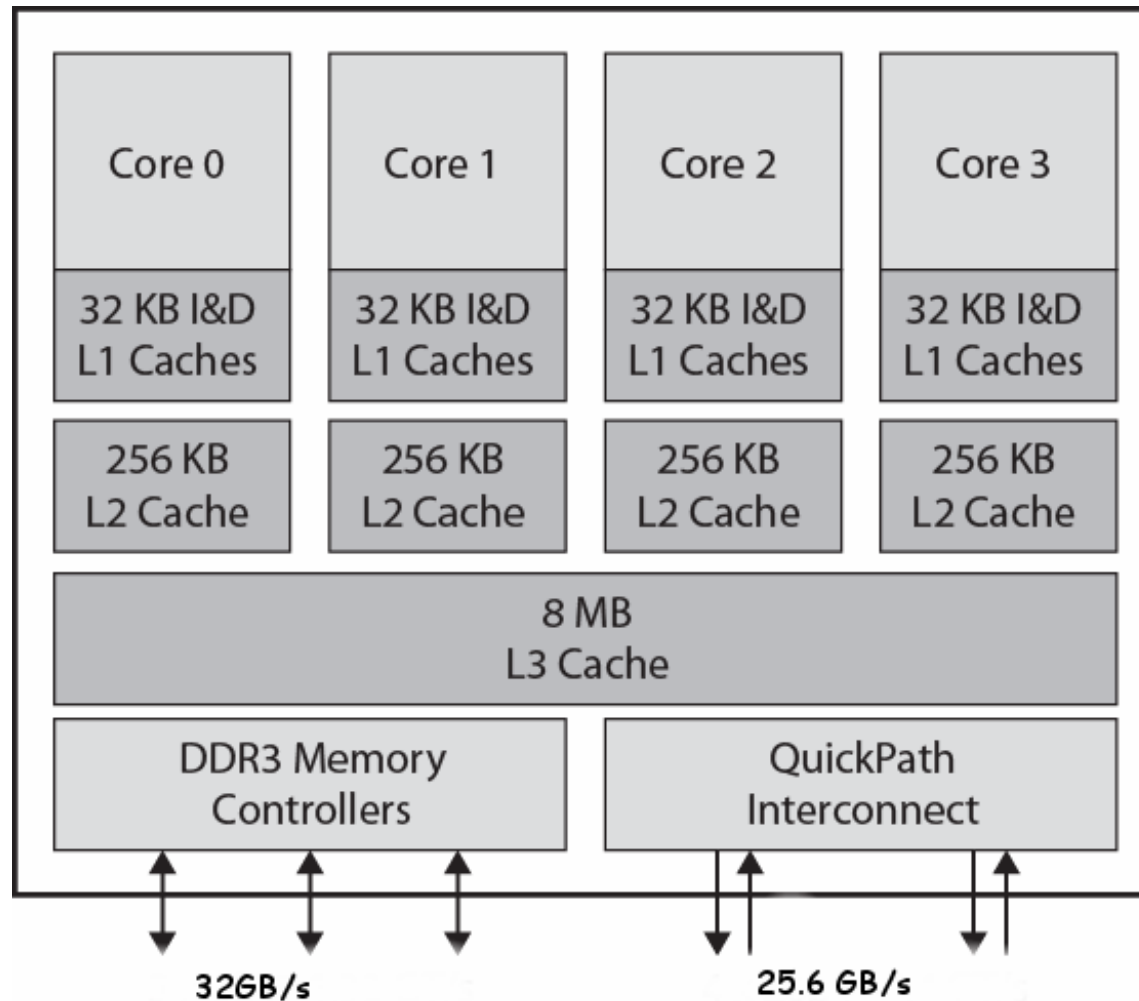


# Intel Core i7 Block Diagram



.3 ns/B !

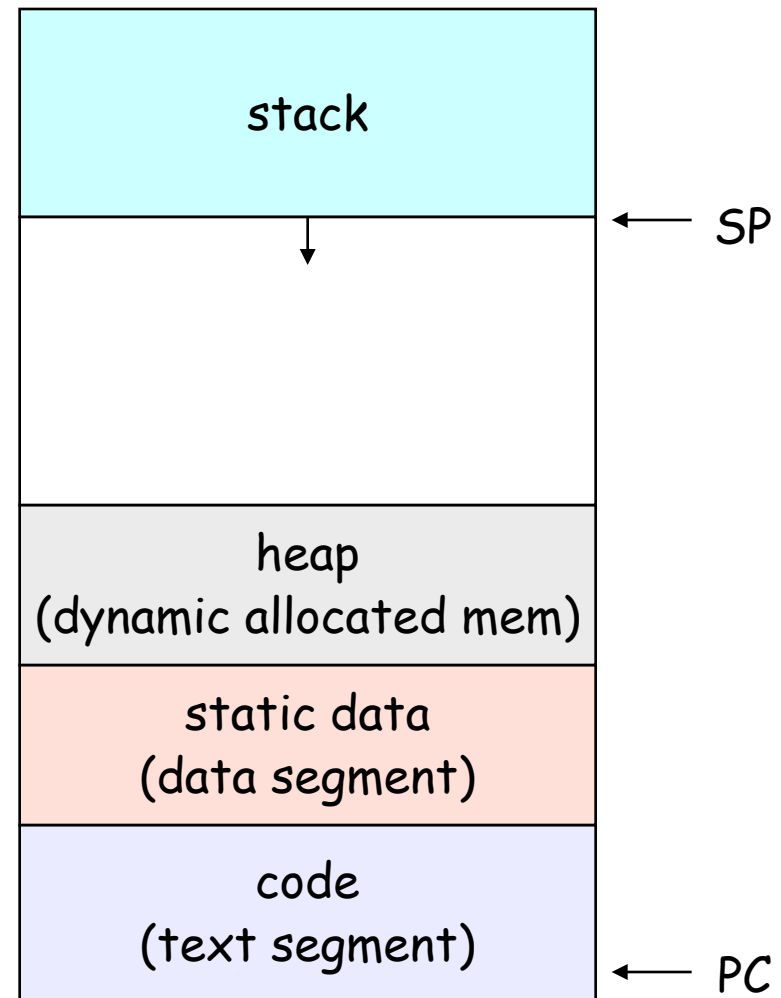
# Intel Core i7 Block Diagram



.3 ns/B !

# Process Address Space Review

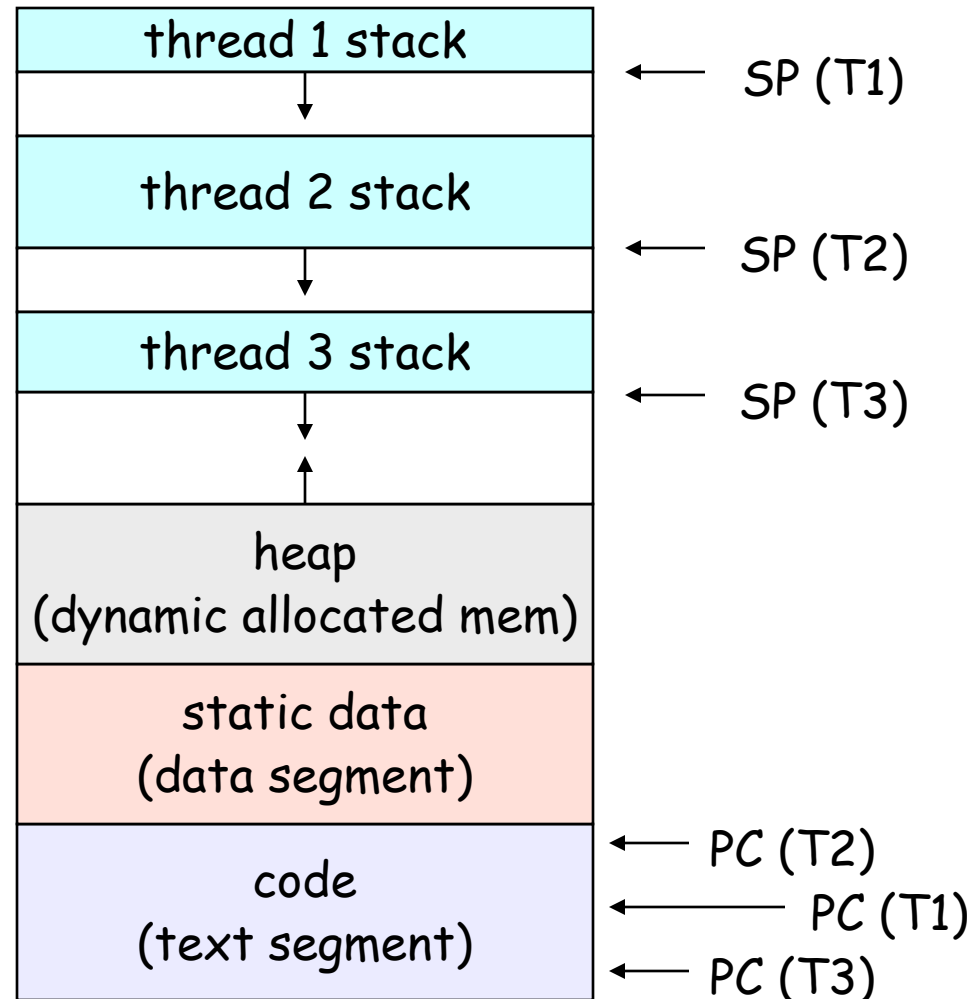
- ❑ Every process has a user stack and a program counter
- ❑ In addition, each process has a kernel stack and program counter  
m (not shown here)



# Threaded Address Space

- Every thread *always* has its own user stack and program counter
  - m For both user, kernel threads
- For user threads, there is only a single kernel stack, program counter, PCB, etc.

*User address space (for both user and kernel threads)*



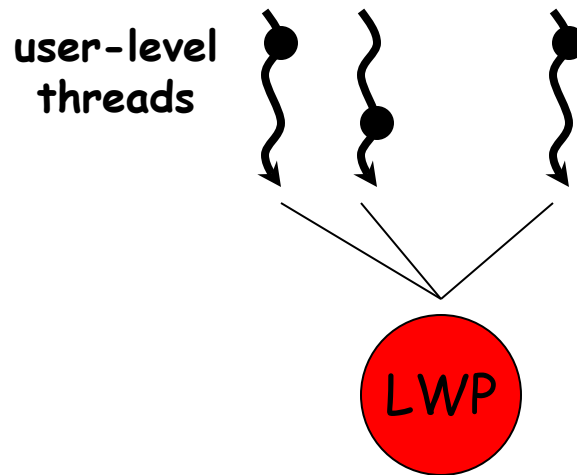
# Light versus Heavy weight

## □ Why Pthreads?

- m The primary motivation for using Pthreads is to realize potential program performance gains.
- m Fork needs IPC (Inter-Process Communication communicate)
- m For example, the following table compares timing results for the `fork()` subroutine and the `pthread_create()` subroutine. Timings reflect 50,000 process/thread creations, were performed with the time utility, and units are in seconds, no optimization flags

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
AMD 2.4 GHz Opteron (8cpus/node)	41.07	60.08	9.01	0.66	0.19	0.43
IBM 1.9 GHz POWER5 p5-575 (8cpus/node)	64.24	30.78	27.68	1.75	0.69	1.10
IBM 1.5 GHz POWER4 (8cpus/node)	104.05	48.64	47.21	2.01	1.00	1.52
INTEL 2.4 GHz Xeon (2 cpus/node)	54.95	1.54	20.78	1.64	0.67	0.90
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.54	1.07	22.22	2.03	1.26	0.67

# Many-to-One Model



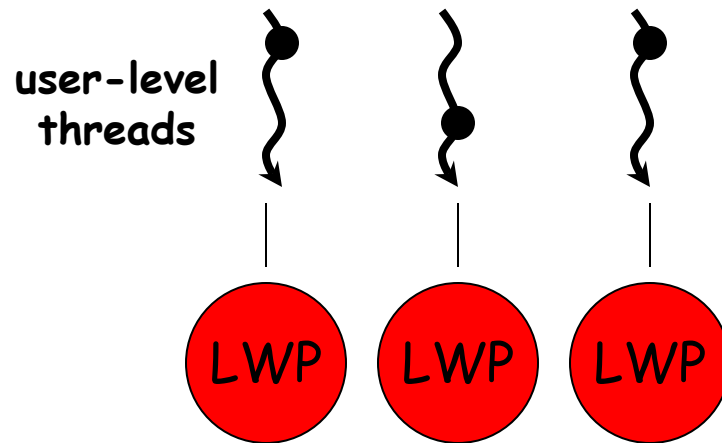
Thread creation, scheduling, synchronization done in user space.

Mainly used in language systems, portable libraries

- ↳ Fast - no system calls required
- ↳ Few system dependencies; portable
- ↳ No parallel execution of threads - can't exploit multiple CPUs
- ↳ All threads block when one uses synchronous I/O



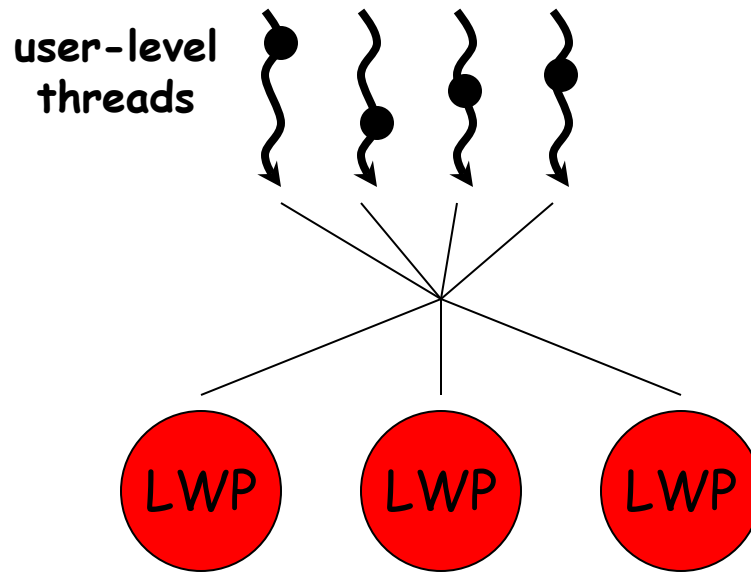
# One-to-one Model



Thread creation, scheduling, synchronization require system calls  
Used in Linux Threads, Windows NT, Windows 2000, OS/2

- More concurrency
- Better multiprocessor performance
- Each user thread requires creation of kernel thread
- Each thread requires kernel resources; limits number of total threads

# Many-to-Many Model



If  $U < L$ ? No benefits of multithreading

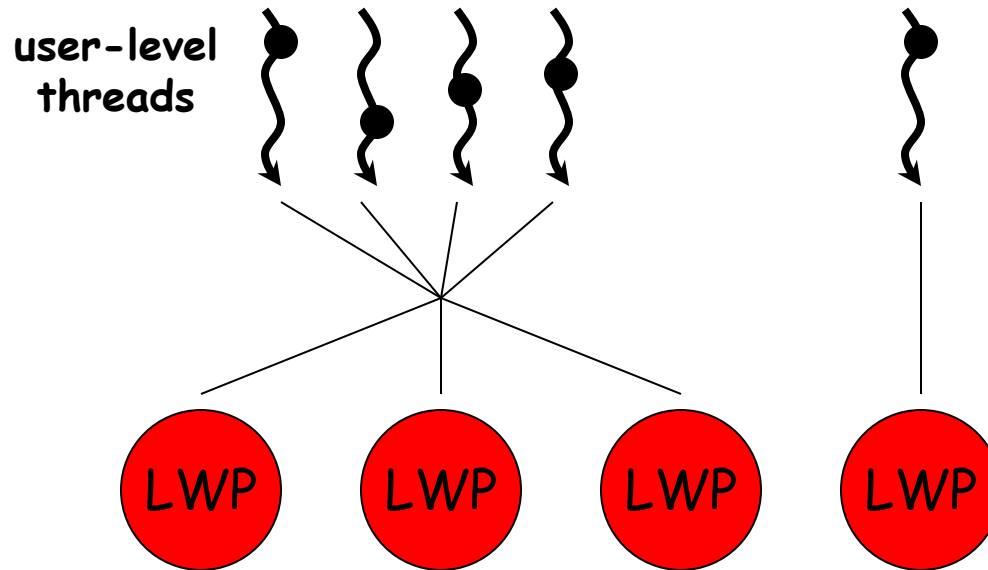
If  $U > L$ , some threads may have to wait for an LWP to run

- Active thread - executing on an LWP
- Runnable thread - waiting for an LWP

A thread gives up control of LWP under the following:

- synchronization, lower priority, yielding, time slicing

# Two-level Model



- Combination of one-to-one + "strict" many-to-many models
- Supports both bound and unbound threads
  - Bound threads - permanently mapped to a single, dedicated LWP
  - Unbound threads - may move among LWPs in set
- Thread creation, scheduling, synchronization done in user space
- Flexible approach, "best of both worlds"
- Used in Solaris implementation of Pthreads and several other Unix implementations (IRIX, HP-UX)

# User-Level vs. Kernel Threads

## User-Level

- ❑ Managed by application
- ❑ Kernel not aware of thread
- ❑ Context switching cheap
- ❑ Create as many as needed
- ❑ Must be used with care

## Kernel-Level

- ❑ Managed by kernel
- ❑ Consumes kernel resources
- ❑ Context switching expensive
- ❑ Number limited by kernel resources
- ❑ Simpler to use

**Key issue:** kernel threads provide virtual processors to user-level threads, but if all of kthreads block, then all user-level threads will block *even if the program logic allows them to proceed*

# When Would User Threads Be Useful?

- ☐ The  $\pi$  calculator?
- ☐ The web server?
- ☐ The Fibonacci GUI?

# Problems with Many-to-Many Threads

- ❑ Lack of coordination between user and kernel schedulers
  - m "Left hand not talking to the right"
- ❑ Specific problems
  - m Poor performance
    - e.g., the OS preempts a thread holding a crucial lock
  - m Deadlock
    - Given  $K$  kernel threads, at most  $K$  user threads can block
      - Other runnable threads are starved out!

# Multi-processor Scheduling

- ❑ Preemptive
- ❑ Non-preemptive