

Lecture 13

□ Administration



Linearizability - formally

- History H is **linearizable** if it can be extended to history G so that G is equivalent to legal sequential history S where $\rightarrow_G \subset \rightarrow_S$.
- G is the same as H but without pending invocations:
 - append responses to pending invocations.
 - discard pending invocations.

Linearizability - formally


Let's explain what is $\rightarrow_G \subset \rightarrow_S$.

Example:

$$\rightarrow_G = \{a \rightarrow c, b \rightarrow c\}$$

$$\rightarrow_S = \{a \rightarrow b, a \rightarrow c, b \rightarrow c\}$$

closure



Example:

Discard
this
pending
invocation:

H =

A q.enq(3)

B q.enq(4)

B q:void


B q.deq()

B q:4

B q.enq(6)

A q:void

Add
response to
this pending
invocation:



Example (cont'):

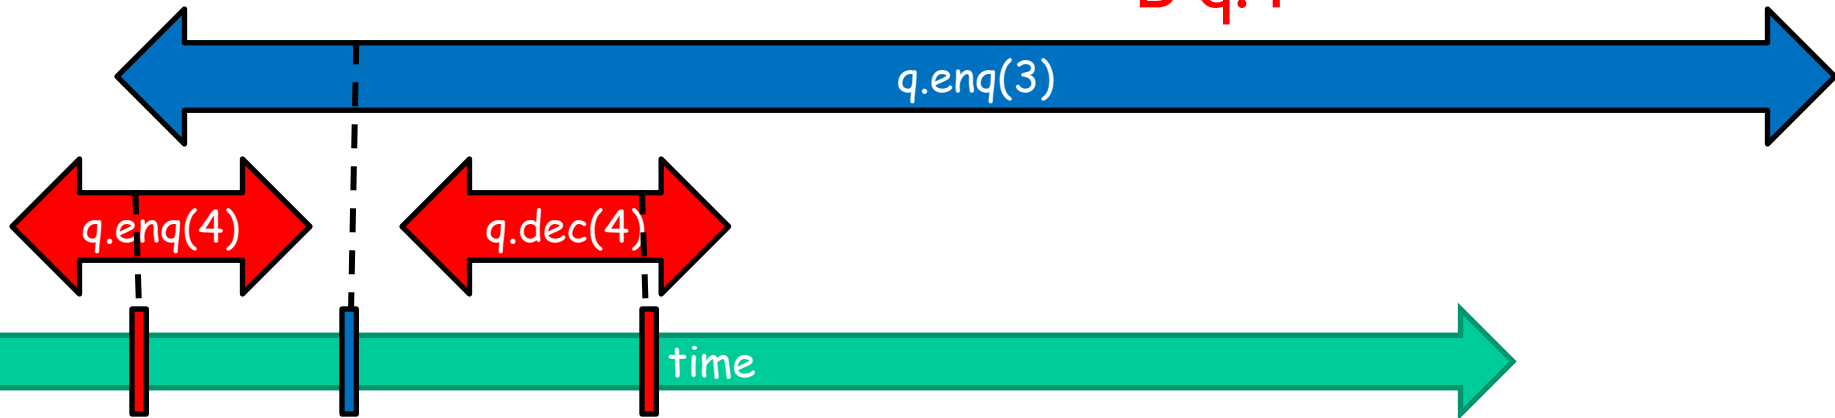
The equivalent sequential history:

G =

- A q.enq(3)
- B q.enq(4)
- B q:void
- B q.deq()
- B q:4
- A q:void

S =

- B q.enq(4)
- B q:void
- A q.enq(3)
- A q:void
- B q.deq()
- B q:4



Composability


- Linearizability also gives us composability:
 - If we want to construct a new object from linearizable objects, we can be sure that our new object is linearizable too.
- why is it good?
 - It gives us modularity. We can prove linearizability independently for each object.

Linearizability: Summary

- ❑ Powerful specification tool for shared objects
- ❑ Allows us to capture the notion of objects being “atomic”
- ❑ Don't leave home without it

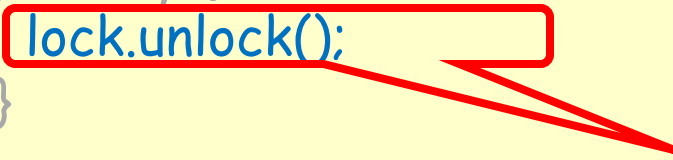
Reasoning About Linearizability: Locking

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```



Reasoning About Linearizability: Locking


```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```



Linearization points
are when locks are
released

More Reasoning: Lock-free

```
public class LockFreeQueue {  
  
    int head = 0, tail = 0;  
    items = (T[]) new Object[capacity];  
  
    public void enq(Item x) {  
        while (tail-head == capacity); // busy-wait  
        items[tail % capacity] = x; tail++;  
    }  
  
    public Item deq() {  
        while (tail == head); // busy-wait  
        Item item = items[head % capacity]; head++;  
        return item;  
    }  
}
```



More Reasoning

Remember that there
is only one enqueuer
and only one dequeuer

```
public class Queue {
    private int head = 0;
    private int tail = 0;
    private Object[] items;

    public void enq(Item x) {
        while (tail - head == capacity); // busy-wait
        items[tail % capacity] = x;
    }

    public Item deq() {
        while (tail == head); // busy-wait
        Item item = items[head % capacity];
        return item;
    }
}
```

Linearization order is
order head and tail
fields modified

tail++;

head++;

Alternative: Sequential Consistency

- History H is *Sequentially Consistent* if it can be extended to G by
 - Appending zero or more responses to pending invocations
 - Discarding other pending invocations
- So that G is equivalent to a
 - Legal sequential history S

Differs from
linearizability

~~Where $\rightarrow_G \subseteq \rightarrow_S$~~



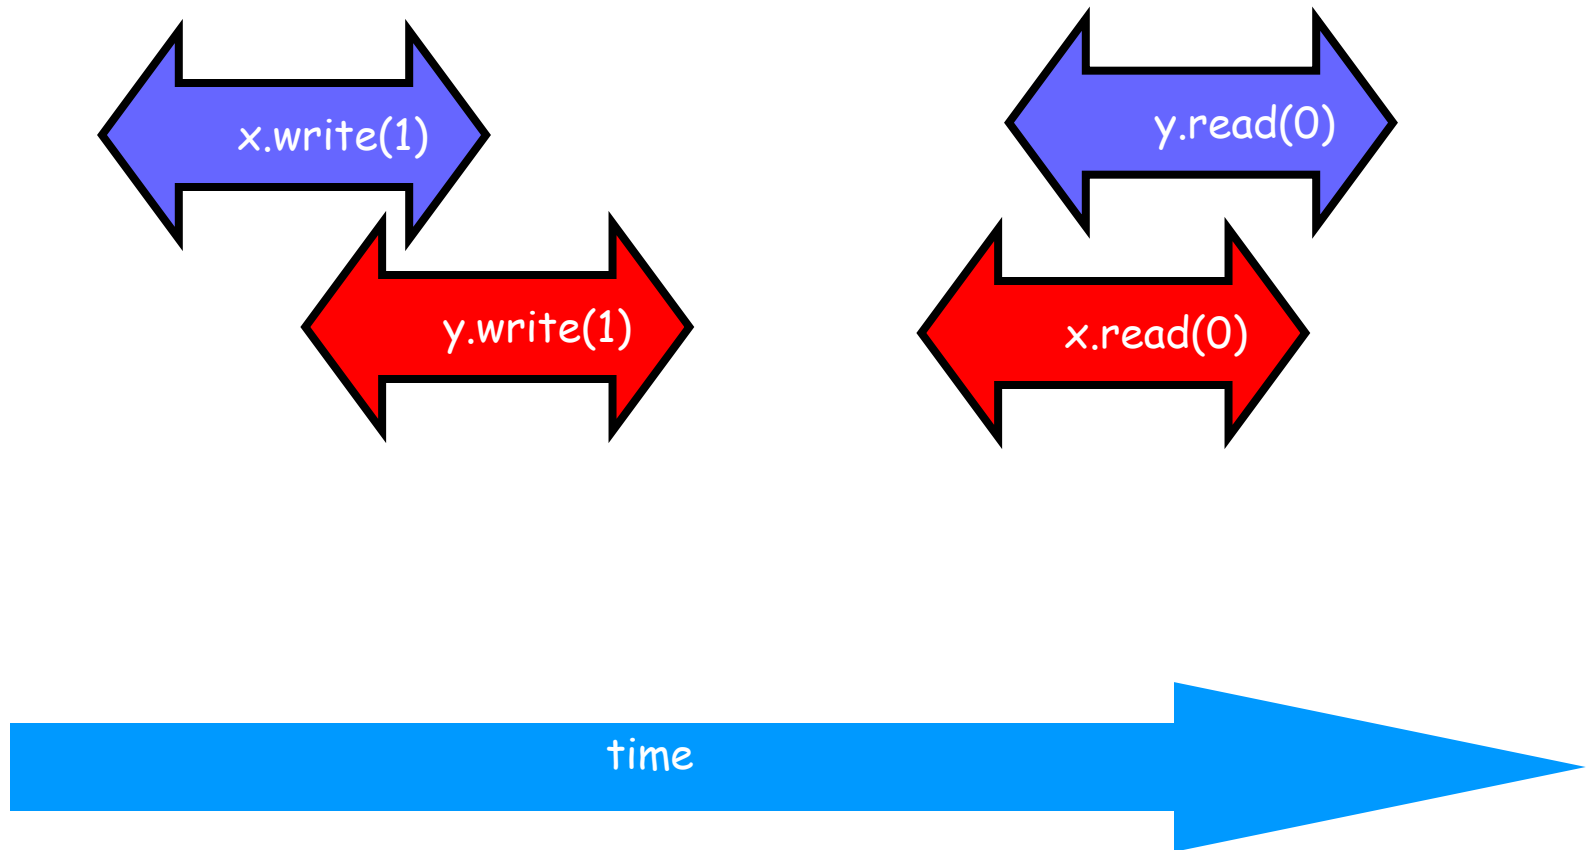
Alternative: Sequential Consistency

- ❑ No need to preserve real-time order
 - m Cannot re-order operations done by the same thread
 - m Can re-order non-overlapping operations done by different threads
- ❑ Often used to describe multiprocessor memory architectures


Fact

- ❑ Most hardware architectures don't support sequential consistency
- ❑ Because they think it's too strong
- ❑ Here's another story ...

The Flag Example




The Flag Example




- Each thread's view is sequentially consistent
 - It went first

The Flag Example



- Entire history isn't sequentially consistent
 - Can't both go first

The Flag Example



- Is this behavior really so wrong?
 - We can argue either way ...

Opinion1: It's Wrong

- ❑ This pattern
 - m Write mine, read yours
- ❑ Is exactly the flag principle
 - m Beloved of Alice and Bob
 - m Heart of mutual exclusion
 - Peterson
 - Bakery, etc.
- ❑ It's non-negotiable!

Opinion2: But It Feels So Right

...

- ❑ Many hardware architects think that sequential consistency (with respect to each data memory location) is too strong
- ❑ Too expensive to implement in modern hardware
- ❑ OK if flag principle
 - m violated by default
 - m Honored by explicit request

Who knew you wanted to synchronize?

- ❑ Writing to memory = mailing a letter
- ❑ Vast majority of reads & writes
 - m Not for synchronization
 - m No need to idle waiting for post office
- ❑ If you want to synchronize
 - m Announce it explicitly
 - m Pay for it only when you need it

Explicit Synchronization

- ❑ Memory barrier instruction
 - m Flush unwritten caches
 - m Bring caches up to date
- ❑ Compilers often do this for you
 - m Entering and leaving critical sections
- ❑ Expensive

Volatile

- ❑ In Java, can ask compiler to keep a variable up-to-date with volatile keyword
- ❑ Also inhibits reordering, removing from loops, & other “optimizations”

Real-World Hardware Memory

- ❑ Weaker than sequential consistency
- ❑ But you can get sequential consistency at a price
- ❑ OK for expert, tricky stuff
 - m assembly language, device drivers, etc.
- ❑ Linearizability more appropriate for high-level software

Critical Sections

- ❑ Easy way to implement linearizability
 - m Take sequential object
 - m Make each method a critical section
- ❑ Problems
 - m Blocking
 - m No concurrency

Progress

- ❑ We saw an implementation whose methods were lock-based (deadlock-free)
- ❑ We saw an implementation whose methods did not use locks (lock-free)
- ❑ How do they relate?

Progress Conditions

- ❑ *Deadlock-free: some thread trying to acquire the lock eventually succeeds.*
- ❑ *Starvation-free: every thread trying to acquire the lock eventually succeeds.*
- ❑ *Lock-free: some thread calling a method eventually returns.*
- ❑ *Wait-free: every thread calling a method eventually returns.*

Progress Conditions

	Non-Blocking	Blocking
Everyone makes progress	Wait-free	Starvation-free
Someone makes progress	Lock-free	Deadlock-free

Transactional Memory

- ❑ Software Transactional Memory
- ❑ Some Intel processors now supporting it.

The Road Ahead

- ❑ Concurrent algorithms pose a great challenge
 - ❑ It is "easy" to write
 - m **Correct** algorithms and let efficiency take care of itself
 - m **Efficient** algorithms and let correctness take care of itself
 - ❑ But very hard to write **correct & efficient** algorithms
- Systematically...**

