

Lecture 16

□ Administration

Logical clocks

Causality is important

Not possible to have a global "physical clock"


Causality used to reason, analyze and prove properties of concurrent systems

Systems depends on ordering only NOT the speed at which things take place

Event counting example

- ❑ Three systems: P_0, P_1, P_2
- ❑ Events a, b, c, \dots
- ❑ Local event counter on each system
- ❑ Systems occasionally communicate

Event counting example



Bad ordering:

$e \rightarrow h$

$f \rightarrow k$

Lamport's algorithm

- ❑ Each message carries a timestamp of the sender's clock
- ❑ When a message arrives:
 - m if receiver's clock < message timestamp
 set system clock to (message timestamp + 1)
 - m else do nothing
- ❑ Clock must be advanced between any two events in the same process

Properties of a clock

A system of clocks and a time domain where for every event e using the clock we can assign a value from the time domain.

$$C : H \rightarrow T$$


- Clocks should be monotonically increasing
- Consistent, $e_i \rightarrow e_j \Rightarrow c(e_i) < c(e_j)$

Strongly consistent $e_i \rightarrow e_j \Leftrightarrow c(e_i) < c(e_j)$


Lamport's Clock

- R1: Internal events $C_i = C_i + d$ ($d > 0$, typically $d=1$)
- R2: Piggyback local logical clock on every message, when received do the following:
 - m $C_i = \max(C_i, C \text{ value in message})$
 - m Execute R1 (updates time)
 - m Deliver the message

Event counting example





Example from Notes




Summary

- ❑ Algorithm needs monotonically increasing software counter
- ❑ Incremented at least when events that need to be timestamped occur
- ❑ Each event has a **Lamport timestamp** attached to it
- ❑ For any two events, where $a \rightarrow b$:
 $L(a) < L(b)$





Problem: Identical timestamps



$a \rightarrow b, b \rightarrow c, \dots$: local events sequenced

$i \rightarrow c, f \rightarrow d, d \rightarrow g, \dots$: Lamport imposes a *send* \rightarrow *receive* relationship

Concurrent events (e.g., a and i) may have the same timestamp ... or not

Unique timestamps (total ordering)

We can force each timestamp to be unique

m Define global logical timestamp (T_i, i)

- T_i represents local Lamport timestamp
- i represents process number (globally unique)
 - E.g. (host address, process ID)

m Compare timestamps:

$$(T_i, i) < (T_j, j)$$


if and only if

$$T_i < T_j \text{ or}$$

$$T_i = T_j \text{ and } i < j$$

Does not relate to event ordering

Unique (totally ordered) timestamps



Problem: Detecting causal relations

If $L(e) < L(e')$

m Cannot conclude that $e \rightarrow e'$

Looking at Lamport timestamps

m Cannot conclude which events are causally related

Solution: use a **vector clock**

Vector clocks

Rules:

1. Vector initialized to 0 at each process
 $V_i[j] = 0$ for $i, j = 1, \dots, N$
2. Process increments its element of the vector in local vector before timestamping event:
 $V_i[i] = V_i[i] + 1$
3. Message is sent from process P_i with V_i attached to it
4. When P_j receives message, compares vectors element by element and sets local vector to higher of two values
 $V_j[i] = \max(V_i[i], V_j[i])$ for $i = 1, \dots, N$

Comparing vector timestamps

Define

$V = V'$ iff $V[i] = V'[i]$ for $i = 1 \dots N$

$V \leq V'$ iff $V[i] \leq V'[i]$ for $i = 1 \dots N$

For any two events e, e'

if $e \rightarrow e'$ then $V(e) < V(e')$


- Just like Lamport's algorithm

if $V(e) < V(e')$ then $e \rightarrow e'$

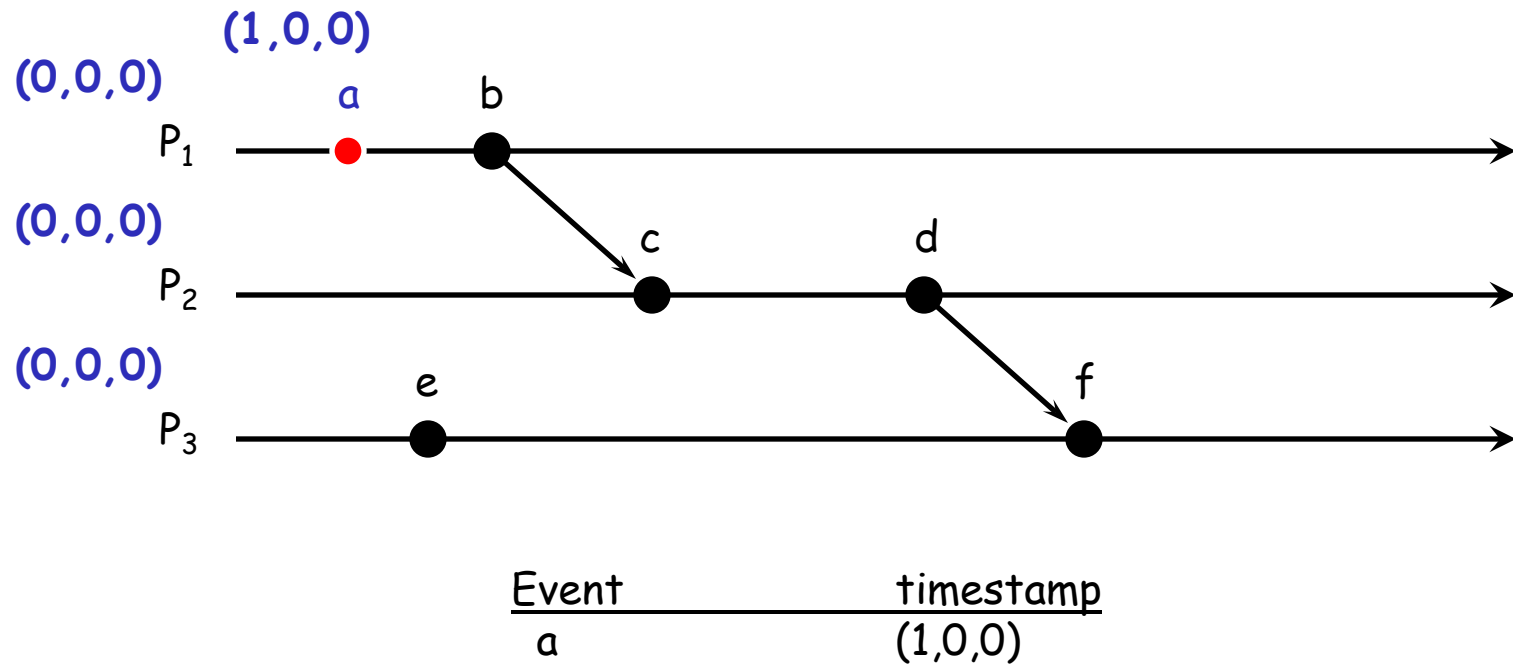
Two events are **concurrent** if neither

$V(e) \leq V(e')$ nor $V(e') \leq V(e)$


Vector timestamps



Vector timestamps




Vector timestamps




<u>Event</u>	<u>timestamp</u>
a	$(1,0,0)$
b	$(2,0,0)$

Vector timestamps




<u>Event</u>	<u>timestamp</u>
a	$(1,0,0)$
b	$(2,0,0)$
c	$(2,1,0)$

Vector timestamps




<u>Event</u>	<u>timestamp</u>
a	$(1,0,0)$
b	$(2,0,0)$
c	$(2,1,0)$
d	$(2,2,0)$

Vector timestamps



<u>Event</u>	<u>timestamp</u>
a	(1,0,0)
b	(2,0,0)
c	(2,1,0)
d	(2,2,0)
e	(0,0,1)
f	(2,2,2)


Vector timestamps



Event	timestamp
a	(1,0,0)
b	(2,0,0)
c	(2,1,0)
d	(2,2,0)
e	(0,0,1)
f	(2,2,2)

concurrent events

Another Example



Vector Clock Assignments

Summary: Logical Clocks & Partial Ordering

□ Causality

m If $a \rightarrow b$ then event a can affect event b

□ Concurrency

m If neither $a \rightarrow b$ nor $b \rightarrow a$ then one event cannot affect the other

□ Partial Ordering

m Causal events are sequenced

□ Total Ordering

m All events are sequenced