

# Lecture 20

On the ugrad machine:

use mpich

Use the package from

`"/cs/local/lib/pkg"`

Use can also download and install  
it yourself from [mpich.org](http://mpich.org)

# MPI is Simple

- Many parallel programs can be written using just these six functions, only two of which are non-trivial:

- m MPI\_Init

- m MPI\_Finalize

- m MPI\_Comm\_size

- m MPI\_Comm\_rank

- m MPI\_Send

- m MPI\_Recv


# Simple Hello (C)

```
#include "mpi.h"
#include <stdio.h>

int main( int argc, char *argv[] )
{
    int rank, size;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );
    printf( "I am %d of %d\n", rank, size );
    MPI_Finalize();
    return 0;
}
```

# MPI Basic Send/Receive

- We need to fill in the details in



- Things that need specifying:
  - m How will “data” be described?
  - m How will processes be identified?
  - m How will the receiver recognize/screen messages?
  - m What will it mean for these operations to complete?

# MPI Basic (Blocking) Send

`MPI_Send (start, count, datatype, dest, tag, comm)`

- ❑ The message buffer is described by (`start`, `count`, `datatype`).
- ❑ The target process is specified by `dest`, which is the rank of the target process in the communicator specified by `comm`.
- ❑ When this function returns, the data has been delivered to the system and the buffer can be reused. The message may not have been received by the target process.

# MPI Basic (Blocking) Receive

`MPI_Recv(start, count, datatype, source, tag, comm, status*)`

- ❑ Waits until a matching (on `source` and `tag`) message is received from the system, and the buffer can be used.
- ❑ `source` is rank in communicator specified by `comm`, or `MPI_ANY_SOURCE`.
- ❑ `status` structure contains further information
- ❑ Receiving fewer than `count` occurrences of `datatype` is OK, but receiving more is an error.

# Identifying Processes

## ❑ MPI Communicator

- m Defines a *group* (set of ordered processes) and a *context* (a virtual network)

## ❑ Rank

- m Process number within the group
- m MPI\_ANY\_SOURCE will receive from any process

## ❑ Default communicator

- m MPI\_COMM\_WORLD the whole group

# Identifying Messages

- ❑ An MPI Communicator defines a virtual network, `send/recv` pairs must use the same communicator
- ❑ `send/recv` routines have a *tag* (integer variable) argument that can be used to identify a message, or screen for a particular message.
  - m `MPI_ANY_TAG` will receive a message with any tag



# Identifying Data

- Data is described by a triple (address, type, count)
  - m For `send`, this defines the message
  - m For `recv`, this defines the size of the receive buffer
- Amount of data received, source, and tag available via *status* data structure
  - m Useful if using `MPI_ANY_SOURCE`, `MPI_ANY_TAG`, or unsure of message size (must be smaller than buffer)

# MPI Types

- Type may be recursively defined as:
  - m An MPI predefined type
  - m A contiguous array of types
  - m An array of equally spaced blocks
  - m An array of arbitrary spaced blocks
  - m Arbitrary structure
- Each user-defined type constructed via an MPI routine, e.g. `MPI_TYPE_VECTOR`

# MPI Predefined Types

C:

MPI\_INT

MPI\_FLOAT

MPI\_DOUBLE

MPI\_CHAR

MPI\_UNSIGNED

MPI\_LONG

Language Independent:

MPI\_BYTE

Fortran:

MPI\_INTEGER

MPI\_REAL

MPI\_DOUBLE\_PRECISION

MPI\_CHARACTER

MPI\_LOGICAL

MPI\_COMPLEX

# MPI Types

- Explicit data description is useful:
  - m Simplifies programming, e.g. send row/column of a matrix with a single call
  - m Heterogeneous machines
  - m May improve performance
    - Reduce memory-to-memory copies
    - Allow use of scatter/gather hardware
  - m May hurt performance
    - User packing of data likely faster

# Point-to-point Example

## Process 0

```
#define TAG 999
float a[10];
int dest=1;
MPI_Send(a, 10,
MPI_FLOAT, dest, TAG,
MPI_COMM_WORLD);
```

## Process 1

```
#define TAG 999
MPI_Status status;
int count;
float b[20];
int sender=0;
MPI_Recv(b, 20,
MPI_FLOAT, sender, TAG,
MPI_COMM_WORLD,
&status);
MPI_Get_count(&status,
MPI_FLOAT, &count);
```

# MPI Basic (Standard) Send

`MPI_Send (start, count, datatype, dest, tag, comm)`

- ❑ The message buffer is described by (`start`, `count`, `datatype`).
- ❑ The target process is specified by `dest`, which is the rank of the target process in the communicator specified by `comm`.
- ❑ When this function returns, the data has been delivered to the system and the buffer can be reused. The message may not have been received by the target process.

# MPI Basic (Standard) Receive

`MPI_Recv(start, count, datatype, source, tag, comm, status*)`

- ❑ Waits until a matching (on `source` and `tag`) message is received from the system, and the buffer can be used.
- ❑ `source` is rank in communicator specified by `comm`, or `MPI_ANY_SOURCE`.
- ❑ `status` structure contains further information
- ❑ Receiving fewer than `count` occurrences of `datatype` is OK, but receiving more is an error.

# MPI Status Data Structure

## □ In C

```
MPI_Status status;  
int recvd_tag, recvd_from, recvd_count;  
recvd_tag = status.MPI_TAG;  
recvd_from = status.MPI_SOURCE;  
MPI_Get_count( &status, MPI_INT,  
               &recvd_count );
```



# MPI's Non-blocking Operations

- ❑ Non-blocking operations return (immediately) "request handles" that can be tested and waited on. (**Posts a send/receive**)

```
MPI_Request request;
```

```
MPI_Isend(start, count, datatype,  
          dest, tag, comm, &request)
```

```
MPI_Irecv(start, count, datatype,  
          dest, tag, comm, &request)
```

```
MPI_Wait(&request, &status)
```

- ❑ One can also test without waiting:

```
MPI_Test(request, &flag, status)
```

# Example

```
#define MYTAG 123
```

```
#define WORLD MPI_COMM_WORLD
```

```
MPI_Request request;
```

```
MPI_Status status;
```

Process 0:

```
MPI_Irecv(B, 100, MPI_DOUBLE, 1, MYTAG, WORLD,  
&request)
```

```
MPI_Send(A, 100, MPI_DOUBLE, 1, MYTAG, WORLD)
```

```
MPI_Wait(&request, &status)
```

Process 1:

```
MPI_Irecv(B, 100, MPI_DOUBLE, 0, MYTAG, WORLD,  
&request)
```

```
MPI_Send(A, 100, MPI_DOUBLE, 0, MYTAG, WORLD)
```

```
MPI_Wait(&request, &status)
```

# Using Non-Blocking Send

Also possible to use non-blocking send:

m "status" argument to **MPI\_Wait** doesn't return useful info here.

```
#define MYTAG 123
```

```
#define WORLD MPI_COMM_WORLD
```

```
MPI_Request request;
```

```
MPI_Status status;
```

```
p=1-me; /* calculates partner in exchange */
```

Process 0 and 1:

```
MPI_Isend(A, 100, MPI_DOUBLE, p, MYTAG, WORLD,  
          &request)
```

```
MPI_Recv(B, 100, MPI_DOUBLE, p, MYTAG, WORLD,  
         &status)
```

```
MPI_Wait(&request, &status)
```

# Non-Blocking Gotchas

## ❑ Obvious caveats:

- m 1. You may not modify the buffer between `Isend()` and the corresponding `Wait()`. Results are undefined.
- m 2. You may not look at or modify the buffer between `Irecv()` and the corresponding `Wait()`. Results are undefined.
- m 3. You may not have two pending `Irecv()`s for the same buffer.

## ❑ Less obvious:

- m 4. You may not *look* at the buffer between `Isend()` and the corresponding `Wait()`.
- m 5. You may not have two pending `Isend()`s for the same buffer.

## ❑ Why the `isend()` restrictions?

- m Restrictions give implementations more freedom, e.g.,
  - Heterogeneous computer with differing byte orders
  - Implementation swap bytes in the original buffer

# Multiple Completions

- It is sometimes desirable to wait on multiple requests:

```
MPI_Waitall(count, array_of_requests,  
            array_of_statuses)
```

```
MPI_Waitany(count, array_of_requests,  
            &index, &status)
```

```
MPI_Waitsome(count, array_of_requests,  
             array_of_indices, array_of_statuses)
```

- There are corresponding versions of test for each of these.

# Multiple completion

- ❑ Source of non-determinism (new issues fairness?), process what is ready first
- ❑ Latency hiding, parallel slack
- ❑ Still need to poll for completion, {do some work; check for comm}\*
  - ❑ **Alternative:** multiple threads or co-routine like support

# MPI point to point routines

❑ MPI_Send	Standard send
❑ MPI_Recv	Standard receive
❑ MPI_Bsend	Buffered send
❑ MPI_Rsend	Ready send
❑ MPI_Ssend	Synchronous send
❑ MPI_Ibsend	Nonblocking, buffered send
❑ MPI_Irecv	Nonblocking receive
❑ MPI_Irsend	Nonblocking, ready send
❑ MPI_Isend	Nonblocking send
❑ MPI_Issend	Nonblocking synchronous send
❑ MPI_Sendrecv	Exchange
❑ MPI_Sendrecv_replace	Exchange, same buffer
❑ MPI_Start	Persistent communication

# Communication Modes

- ❑ Standard
  - m Usual case (system decides)
  - m **MPI\_Send, MPI\_Isend**
- ❑ Synchronous
  - m The operation does not complete until a matching receive has started copying data into its receive buffer. (no buffers)
  - m **MPI\_Ssend, MPI\_Issend**
- ❑ Ready
  - m Matching receive already posted. (0-copy)
  - m **MPI\_Rsend, MPI\_Irsend**
- ❑ Buffered
  - m Completes after being copied into user provided buffers (Buffer\_attach, Buffer\_detach calls)
  - m **MPI\_Bsend, MPI\_Ibsend**

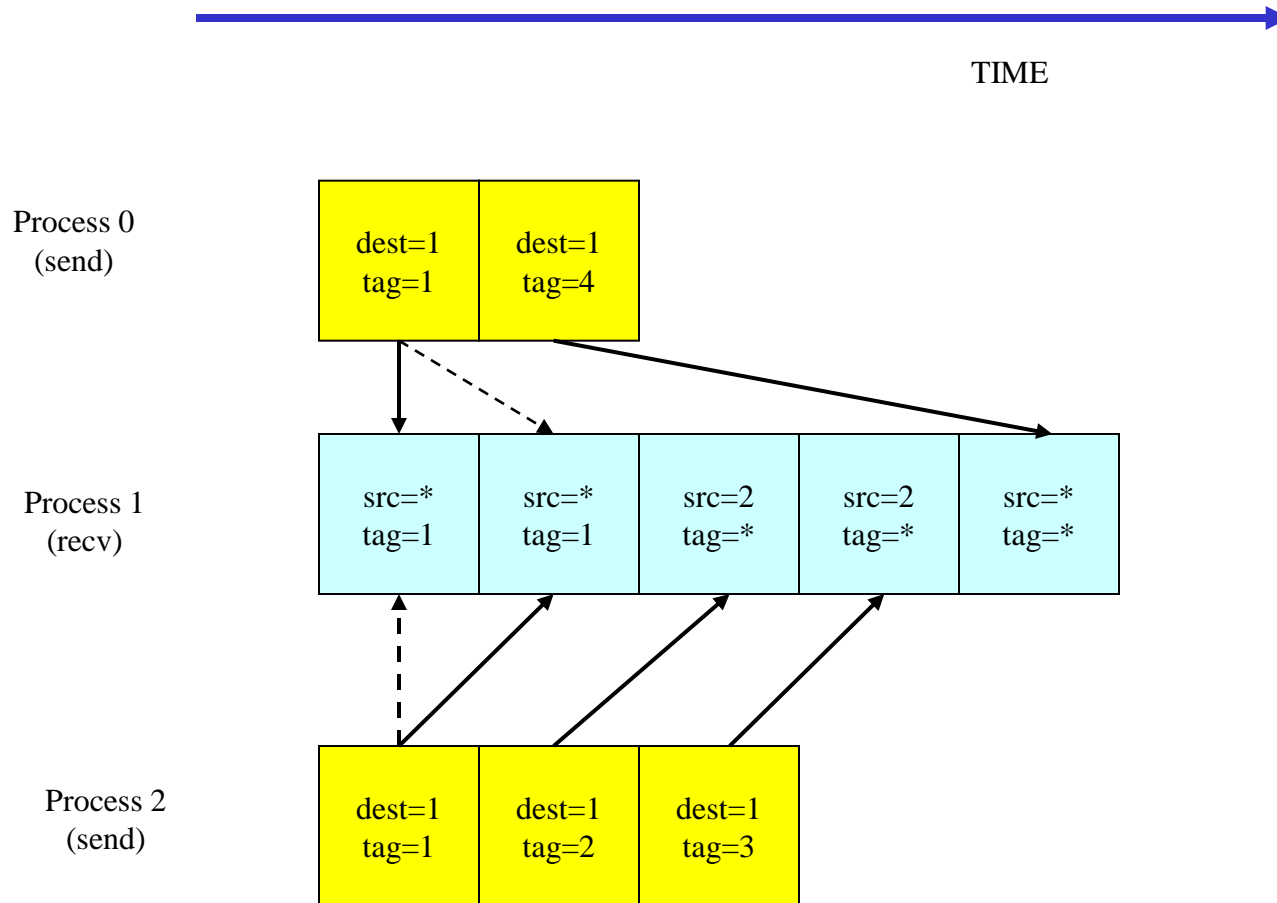


# Point to point with modes

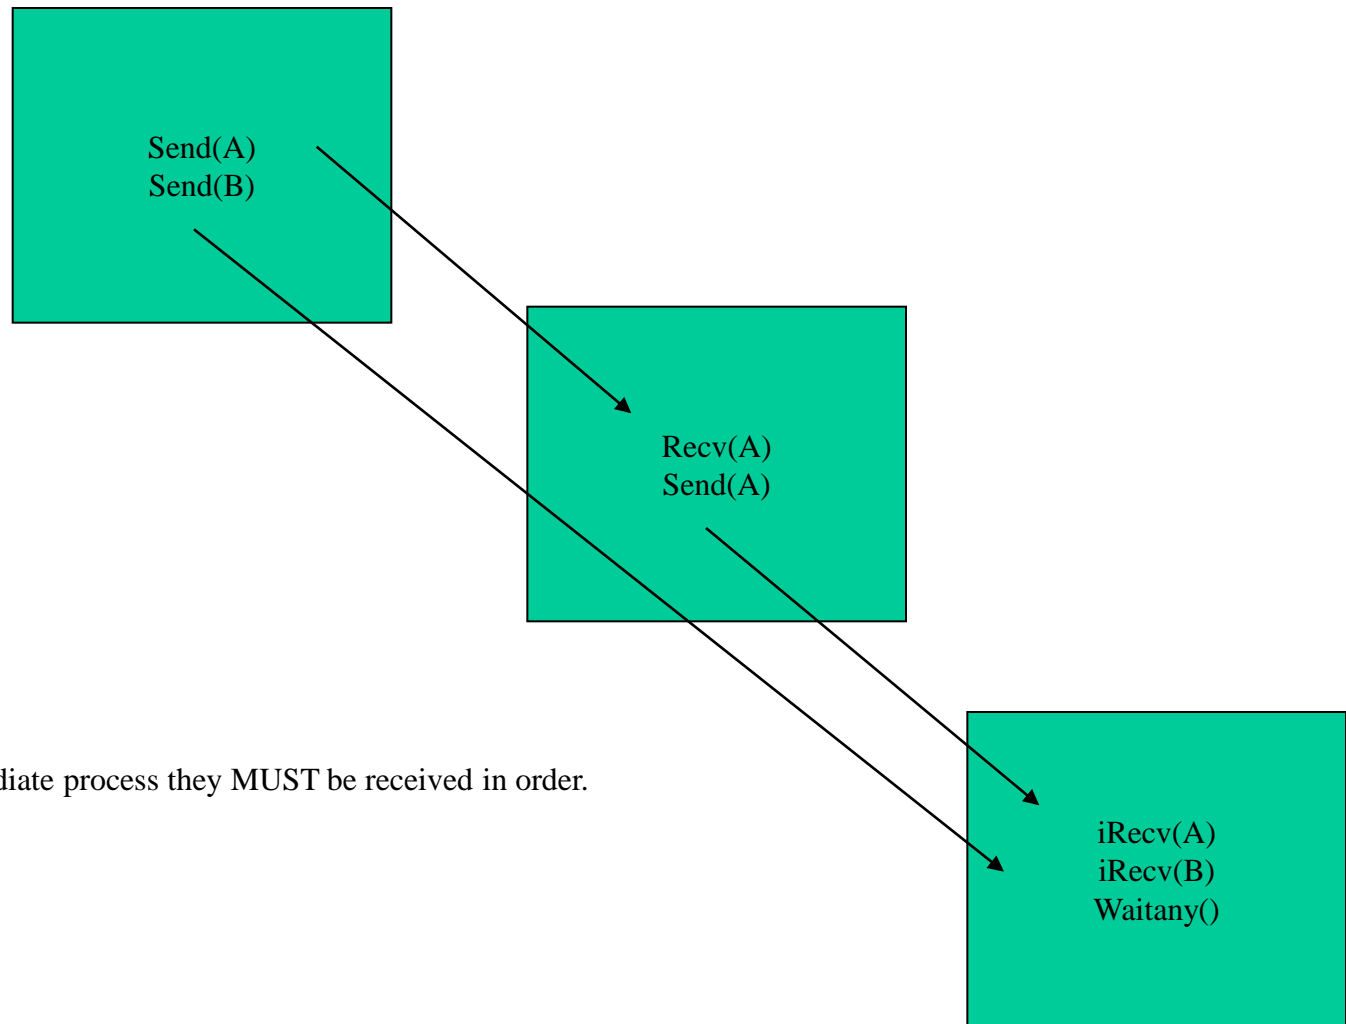
`MPI_[SBR]send(start, count, datatype, dest, tag, comm)`

There is only one mode for receive!

# Messages matched in order



# Message ordering



Without the intermediate process they **MUST** be received in order.

# Sources of Deadlocks

- ❑ Send a large message from process 0 to process 1
  - m If there is insufficient storage at the destination, the send must wait for the user to provide the memory space (through a receive)
- ❑ What happens with this code?

Process 0

Process 1

---

**Send( 1 )**

**Send( 0 )**

**Recv( 1 )**

**Recv( 0 )**

- This is called “unsafe” because it depends on the availability of system buffers

# Some Solutions to the "unsafe" Problem

- Order the operations more carefully:

Process 0	Process 1
<b>Send( 1 )</b>	<b>Recv( 0 )</b>
<b>Recv( 1 )</b>	<b>Send( 0 )</b>

Supply receive buffer at same time as send:

Process 0	Process 1
<b>Sendrecv( 1 )</b>	<b>Sendrecv( 0 )</b>

# More Solutions to the "unsafe" Problem

- ❑ Supply own space as buffer for send

Process 0

Process 1

---

**Bsend(1)**

**Bsend(0)**

**Recv(1)**

**Recv(0)**

Use non-blocking operations:

Process 0

Process 1

---

**Irecv(1)**

**Irecv(0)**

**Isend(1)**

**Isend(0)**

**Waitall**

**Waitall**