

Lecture 21

□ Outline

- m Leader Election

- m Complexity measure of distributed computation

- Bits
- Messages
- Rounds

- m Anonymous Ring - leader election

- Symmetry breaking
- Coin tossing

- m Ring with IDs, no faults - leader election

- Chang - Roberts Algorithm with Participants/Non-Participants

Motivation

- We often need a coordinator in distributed systems
 - m Leader, distinguished node/process


- If we have a leader, mutual exclusion is trivially solved
 - m The leader determined who enters CS
 - m Inefficient though...

Election Algorithms


- ❑ Any process can serve as coordinator
- ❑ Any process can "call an election" (initiate the algorithm to choose a new coordinator).
 - m There is no harm (other than extra message traffic) in having multiple concurrent elections.
- ❑ Elections may be needed when the system is initialized, or if the coordinator crashes or retires.

Leader Election in Ring Networks

Initial state
(all not-elected)




Final state



Sense-of-direction in Rings

- In an *oriented* ring, processors have a consistent notion of left and right



1 = left = clockwise

2 = right = counter-clockwise

- For example, if messages are always forwarded on channel 1, they will cycle clockwise around the ring

Properties of Rings and Algorithms

Anonymous Ring (no identifiers)

Non-anonymous Ring


Size of the network n is known (non-uniform)

Size of the network n is not known (uniform)

Synchronous Algorithm

Asynchronous Algorithm

Leader Election in Anonymous Rings



Every processor runs the same algorithm

Every processor does exactly the same execution

Impossibility for Anonymous Rings

- ❑ **Theorem:** There is *no* leader election algorithm for anonymous rings, even if
 - the algorithm knows the ring size (non-uniform)
 - in the synchronous model
- ❑ **Proof Sketch (for non-uniform and synchronous rings):**
 - m Every processor begins in same state (*not-elected*) with same outgoing msgs (since anonymous)
 - m Every processor receives same msgs, does same state transition, and sends same msgs in round 1
 - m And so on and so forth for rounds 2, 3, ...
 - m Eventually some processor is supposed to enter an elected state. But then they all would.

Definition of an Election Algorithm

- Formally, an algorithm is an election algorithm if and only if:
 - m Each process has the same local algorithm
 - m The algorithm is decentralized
 - It can be initiated by any number of processes in the system
 - m It reaches a terminal configuration, and in each reachable terminal configuration one process is in state *leader* and the rest are in the state *lost*

Assumptions

- ❑ Fully asynchronous system
 - m No global clock, transmission times arbitrary

- ❑ Every process has a unique identity
 - m Identities are totally ordered
 - m Because we have finite number of processes:
 - *Max* identity and *Min* identity available (extreme values)

LeLann's ring election

- Whenever a node receives back its id, it has seen every other initiators id
 - m Assuming FIFO channels

- Let every node keep a list of every identifier seen ($list_p$)
 - m If non-initiator, $state=lost$ immediately
 - m If initiator, when own id received:
 - $state=leader$ if $\min\{list_p\}=p$
 - $state=lost$ otherwise

LeLann's Algorithm

var $List_p$: set of \mathcal{P} **init** $\{p\}$; *Initially only know myself*
 $state_p$;

begin *if* p is initiator **then**

Send my id, and wait

*Repeat forwarding
and collecting ids
until we receive
our id*

Termination:

did we win or lose

begin $state_p := cand$; **send** $\langle tok, p \rangle$ to $Next_p$; **receive** $\langle tok, q \rangle$;

while $q \neq p$ **do**

begin $List_p := List_p \cup \{q\}$;

send $\langle tok, q \rangle$ to $Next_p$; **receive** $\langle tok, q \rangle$

end ;

if $p = \min(List_p)$ **then** $state_p := leader$
else $state_p := lost$

end

else while *true* **do**

*Non-initiators just
forward and lose*

begin **receive** $\langle tok, q \rangle$; **send** $\langle tok, q \rangle$ to $Next_p$;

if $state_p = sleep$ **then** $state_p := lost$

end

end