

Lecture 23

□ Administration

Chapter 4: *Global State and snapshot recording algorithms*




Bully Algorithm

1. Send *election* message (*I want to be the leader*) to processes with *larger id*
2. Give up your bid if a process with *larger id* sends a *reply* message (*means no, you cannot be the leader*). In that case, wait for the *leader* message (*I am the leader*). Otherwise elect yourself the leader and send a *leader* message
3. If *no reply is received*, then elect yourself the leader, and broadcast a *leader* message.
4. If you receive a reply, but later don't receive a *leader* message from a process of larger id (i.e the leader-elect has crashed), then re-initiate election by sending *election* message.






The process q now calls an election (if it has not already done so).

Repeat until no higher-level process responds. The last process to call an election "wins" the election.

The winner sends a message to other processes announcing itself as the new coordinator.




If 7 comes back on line, it will call an election




The Bully Algorithm


The coordinator p_4 fails and p_1 detects this starts election



On the message from p_1 , p_2 and p_3 start their own election




Before p_3 can announce victory it fails, assuming p_1 timeouts first.




Eventually.....

Eventually p_2 can announce victory .



Worst Case



Node 0 sends N-1 **election** messages
over again

Node 1 sends N-2 **election** messages

Node N-2 sends 1 **election** messages etc

So, 0 starts all

Finally, node N-2 will be elected leader, but
before it sent the **leader** message, it crashed.

The worst-case message complexity = $O(n^3)$ (This is bad)


Synchronizers

Synchronous algorithms (round-based, where processes execute actions in lock-step synchrony) are easier to deal with than ***asynchronous algorithms***. In each round, a process

- (1) receives messages from neighbors,
- (2) performs local computation
- (3) sends messages to ≥ 0 neighbors

A synchronizer is a protocol that enables synchronous algorithms to run on asynchronous platforms

Synchronous algorithm



Asynchronous system

Synchronizers

Simulate a synchronous network over an asynchronous underlying network

Possible in the absence of failures

Enables us to use simple synchronous algorithms even when the underlying network is asynchronous

Synchronous network abstraction: A message sent in *pulse* i is received at pulse $i+1$

Synchronizer indicates when a process can generate a pulse


A process can go from pulse i to $i+1$ only when it has received and acted on all messages sent during pulse $i-1$

Synchronizers

In each pulse:

- m A process receives messages sent during the previous pulse
- m It then performs internal computation and sends out messages if required
- m It can execute the next pulse only when the synchronizer permits it


Tree synchronizer



Form a spanning tree with any node as the root. The root initiates the simulation of each tick by sending message $m(j)$ for each clock tick j down the tree. Each process responds with $ack(j)$ and then with a $safe(j)$ message (that represents the fact that the entire subtree under it is safe). When the root receives $safe(j)$ from every child, it initiates the simulation of clock tick $(j+1)$

*Message complexity $M(\beta) = 3(N-1)$
since three messages (m , ack , $safe$) flow along each edge of the tree.*

*Time complexity $T(\beta) = \text{depth of the tree}$.
For a balanced tree, this is $O(\log N)$*



State of Channel

All messages that have been sent but not yet received.

$$S_{ij}^{x,y} = \{m_{ij} : \text{send}(m_{ij}) \leq \text{recv}(m_{i,j}) > LS_j^y\}$$

LS_j^y The state of process j after the occurrence of event e_j^y

Global State

$$GS = \left\{ \bigcup_i LS_i^{x_i}, \bigcup_{j,k} S_{jk}^{y_j, z_k} \right\}$$

Consistent or Inconsistent

Terms

- ❑ Concurrent
- ❑ Cut some Global State
- ❑ Consistent, transitless (no outstanding messages), strongly consistent (consistent and transitless)

Chandy and Lamport Snapshot

Marker-Sending Rule for a Process p :

p records its state;

for (each channel C directed away from p , with a marker not sent)
 { p sends one marker along C before p sends any further
 messages along C ; }

Marker-Receiving Rule for a Process q :


if (q has not recorded its state) then

 { q records the state of C as the empty sequence;
 execute marker-sending rule.
 }

else { q records the state of C as the sequence of message
 received along C after q 's state was recorded and before
 q received the marker along C .
 }


- ❑ When a process receives a mark, it knows that a snapshot is in process.
- ❑ An individual node knows that it is done when it records its own state and all the states in my incoming channels.

Example -- initial




Node	Recorded state				
		c1	c2	c3	c4
p			{ }		{ }
q		{ }			
r				{ }	

Example - step 1




Node	Recorded state				
	state	c1	c2	c3	c4
p	490		{}		{}
q		{}			
r				{}	

Example - step 2




Node	Recorded state				
	state	c1	c2	c3	c4
p	490		{}		{}
q	480	{empty}			
r				{}	

Example - step 3




Node	Recorded state				
	state	c1	c2	c3	c4
p	490		{}		{}
q	480	{empty}			
r	485			{empty}	

Example - step 4



Node	Recorded state				
	state	c1	c2	c3	c4
p	490		{20}		{}
q	480	{empty}			
r	485			{empty}	


Example - step 5



Node	Recorded state				
	state	c1	c2	c3	c4
p	490		{20}		{25}
q	480	{empty}			
r	485			{empty}	


What if more than one initiate?

Example -- initial





Node	Recorded state				
		c1	c2	c3	c4
p			{ }		{ }
q		{ }			
r				{ }	


Example - step 5





Node	Recorded state				
	state	c1	c2	c3	c4
p	490		{20}		{25}
q	480	{empty}			
r	485			{empty}	














What if more than one initiate?