

Lecture 24

□ Administration

Chapter 4: *Global State and snapshot recording algorithms*

Chandy and Lamport Snapshot

Marker-Sending Rule for a Process p :

p records its state;

for (each channel C directed away from p , with a marker not sent)
 { p sends one marker along C before p sends any further
 messages along C ; }

Marker-Receiving Rule for a Process q :


if (q has not recorded its state) then

 { q records the state of C as the empty sequence;
 execute marker-sending rule.
 }

else { q records the state of C as the sequence of message
 received along C after q 's state was recorded and before
 q received the marker along C .
 }


- ❑ When a process receives a mark, it knows that a snapshot is in process.
- ❑ An individual node knows that it is done when it records its own state and all the states in my incoming channels.

Example -- initial





Node	Recorded state				
		c1	c2	c3	c4
p			{ }		{ }
q		{ }			
r				{ }	


Example - step 5





Node	Recorded state				
	state	c1	c2	c3	c4
p	490		{20}		{25}
q	480	{empty}			
r	485			{empty}	














What if more than one initiate?

Table 4.1 A comparison of snapshot algorithms.

Algorithms	Features
Chandy–Lamport [7]	Baseline algorithm. Requires FIFO channels. $O(e)$ messages to record snapshot and $O(d)$ time.
Spezialetti–Kearns [29]	Improvements over [7]: supports concurrent initiators, efficient assembly and distribution of a snapshot. Assumes bidirectional channels. $O(e)$ messages to record, $O(rn^2)$ messages to assemble and distribute snapshot.
Venkatesan [32]	Based on [7]. Selective sending of markers. Provides message-optimal incremental snapshots. $\Omega(n + u)$ messages to record snapshot.
Helary [12]	Based on [7]. Uses wave synchronization. Evaluates function over recorded global state. Adaptable to non-FIFO systems but requires inhibition.
Lai–Yang [18]	Works for non-FIFO channels. Markers piggybacked on computation messages. Message history required to compute channel states.
Li <i>et al.</i> [20]	Similar to [18]. Small message history needed as channel states are computed incrementally.
Mattern [23]	Similar to [18]. No message history required. Termination detection (e.g., a message counter per channel) required to compute channel states.
Acharya–Badrinath [1]	Requires causal delivery support. Centralized computation of channel states. Channel message contents need not be known. Requires $2n$ messages, 2 time units.
Alagar–Venkatesan [2]	Requires causal delivery support. Distributed computation of channel states. Requires $3n$ messages, 3 time units, small messages.

$n = \#$ processes, $u = \#$ edges on which messages were sent after previous snapshot, $e = \#$ channels, $d =$ diameter of the network, $r = \#$ concurrent initiators.