


# Lecture 28


□ Administration

# Multicast communication

- ❑ Send message over a distribution tree.
- ❑ Use network hardware support for broadcast or multicast when it is available.
- ❑ Minimize the time and bandwidth utilization




multicast via unicast




network multicast

# Replicated Data Storage







# A network partition




# Reliable Multicast


- Multicast - general idea
- Support
  - m None
  - m IP-multicast




a (allowed).




b (allowed).



c (disallowed).



d (disallowed).





# Reliable Multicast Communication

- ❑ IP multicast uses UDP which means that it is not reliable.
- ❑ An IP multicast message may be lost part way and delivered to some, but not all, of the intended receivers.
- ❑ A process may still send a message to a group using TCP.
  - m This is point to point i.e., this means that the process sends a message to the first replica, then sends a message to the second replica etc;
  - m This is still not reliable. Consider what happens if the sender fails after sending to a subset of the group.
- ❑ Reliable multicast means that every message should be delivered to each current group member

# Reliable Multicast

- ❑ Simple multicasting is sending a message to every process that is a member of a defined group.  
Reliable multicasting requires these properties:
- ❑ **Integrity**—a correct process sends a message to only a member of the group and does it only once.
- ❑ **Validity**—if a correct process sends a message, it will eventually be delivered.
- ❑ **Agreement**—if a message is delivered to a correct process, all other correct processes in the group will deliver it.

# Reliability

*Correct* processes: those that never fail.

## □ Integrity

*A correct process delivers a message at most once.*

## □ Validity

*A message from a correct process will be delivered by the process eventually.*

## □ Agreement

*A message delivered by a correct process will be delivered by all other correct processes in the group.*

⇒ Validity + Agreement = Liveness

# B-multicast

## □ Assumption:

m Reliable one-to-one *send* operation (e.g. TCP)

## □ Basic multicast

m Requirement:

- All correct processes will eventually deliver the message from the *correct* multicaster.

m Implementation:

- *B-multicast*(  $g, m$  ):  $\forall p \in g: \text{send}(p, m)$ ;
  - On receive(  $m$  ) at  $p$ : *B-deliver*(  $m$  ) at  $p$ .
- ⇒ Properties: integrity, validity.

# R-multicast

## □ Reliable multicast

m Requirements: integrity, validity, *agreement*

m Implementation:

- *Received* := {};
- *R-multicast*( *g*, *m*) at process *p*: *B-multicast*( *g*, *m*);
- On *B-deliver*( *m*) at process *q*  
if( *m* ∉ *Received*)  
    *Received* := *Received* ∪ {*m*};  
    if( *q* ≠ *p*) *B-multicast*( *g*, *m*);  
    *R-deliver*( *m*);  
end if

⇒ Inefficient: each message is sent  $|g|$  times to each process

# Reliable Multicast Algorithm

When a message is delivered, the receiving process multicasts it. Duplicate messages are identified (possible by a sequence number) and not delivered.

*On initialization*

*Received* := {};

*For process p to R-multicast message m to group g*

*B-multicast(g, m);      // p ∈ g is included as a destination*

*On B-deliver(m) at process q with g = group(m)*

*if (m ∉ Received)*

*then*

*Received := Received ∪ {m};*

*if (q ≠ p) then B-multicast(g, m); end if*


*R-deliver m;*

*end if*

# Hold-back queue

A holdback queue for arriving multicast messages that enables the receiving process to obtain metadata about an arriving messages simplifies the implementation of reliable multicast.



Incoming  
messages



# Ordered messages


- ❑ If it is important that messages be delivered in order, there are three types of **ordering**:
- ❑ **FIFO**—(First-in, first-out) if a correct process delivers a message before another, every correct process will deliver the first message before the other.
- ❑ **Casual**—any correct process that delivers the second message will deliver the previous message first.
- ❑ **Total**—if a correct process delivers a message before another, any other correct process that delivers the second message will deliver the first message first.





Not CO -- Processes see different view

CO- if  $m$  sent before  $m_1$




Not CO -- even when the update is consistent

# Comments on Ordering

- Note that **FIFO ordering** and **casual ordering** are only **partial orders**. Not all messages are sent by the same sending process. In addition some multicasts are **concurrent**, not able to be ordered by **happened-before**.
- Note that  $T_1$  and  $T_2$  are delivered in opposite order to the physical time of message creation. Total order demands consistency, but not a particular order.

# Total, FIFO and causal ordering of multicast messages

Notice the consistent ordering of totally ordered messages  $T_1$  and  $T_2$ , the FIFO-related messages  $F_1$  and  $F_2$  and the causally related messages  $C_1$  and  $C_3$  – and the otherwise arbitrary delivery ordering of messages.



# FIFO-ordered Multicast

## □ FIFO-ordered multicast:

### m Assumption:

- a process belongs to at most one group.

### m Implementation:

- Local variables at  $p$ :  $S_p = 1$ ,  $R_p[|g|] = \{0\}$ ;
- ***FO-multicast( $g, m$ )*** at  $p$ :  
    *B-multicast( $g, \langle m, S_p \rangle$ )*;  
     $S_p++$ ;
- On *B-deliver( $\langle m, S \rangle$ )* from  $q$ :  
    if(  $S = R_p[q] + 1$  )  
        ***FO-deliver( $m$ )***;  
         $R_p[q] := S$ ;  
    else if(  $S > R_p[q] + 1$  )  
        place  $\langle m, S \rangle$  in the queue until  $S = R_p[q] + 1$ ;  
        ***FO-deliver( $m$ )***;  
         $R_p[q] := S$ ;  
    end if

# Total Communication

| item | From         | Subject          |
|------|--------------|------------------|
| 23   | A. Hanlon    | Mach             |
| 24   | G. Joseph    | Microkernels     |
| 25   | A. Hanlon    | Re: Microkernels |
| 26   | T.L. Heureux | RPC performance  |
| 27   | M. Walker    | Re: Mach         |