


Lecture 29

□ Administration

R-multicast



□ Reliable multicast

m Requirements: integrity, validity, *agreement*

m Implementation:

- *Received* := {};
- *R-multicast*(*g*, *m*) at process *p*: *B-multicast*(*g*, *m*);
- On *B-deliver*(*m*) at process *q*
if(*m* ∉ *Received*)
 Received := *Received* ∪ {*m*};
 if(*q* ≠ *p*) *B-multicast*(*g*, *m*);
 R-deliver(*m*);
end if

⇒ Inefficient: each message is sent $|g|$ times to each process

Reliability

Correct processes: those that never fail.

❑ Integrity

A correct process delivers a message at most once.

❑ Validity

A message from a correct process will be delivered by the process eventually.

❑ Agreement

A message delivered by a correct process will be delivered by all other correct processes in the group.

⇒ Validity + Agreement = Liveness

Ordered Reliable Multicasts

Assumptions: a process belongs to at most one group.

□ FIFO




m if $m_p \rightarrow m'_p$, all correct processes *that deliver m'_p* will deliver m_p before m'_p .

□ Causal

m if $m \rightarrow m'$, all correct processes *that deliver m'* will deliver m before m' .

□ Total

m if a correct process delivers m before m' , all other correct processes *that deliver m'* will deliver m before m' .



Which multicast semantics? - Why?

item	From	Subject
23	A. Hanlon	Mach
24	G. Joseph	Microkernels
25	A. Hanlon	Re: Microkernels
26	T.L. Heureux	RPC performance
27	M. Walker	Re: Mach

Total Order Specification

□ Validity

If a correct process TO-broadcasts a message m , then it eventually TO-delivers m .

□ Uniform Agreement

If a process TO-delivers a message m , then all correct processes eventually TO-deliver m .

□ Uniform Integrity

For any message m , every process TO-delivers m at most once, and only if m was previously TO-broadcast by $\text{sender}(m)$.

□ Uniform Total Order

If processes p and q both TO-deliver messages m and m' , then p TO-delivers m before m' , if and only if q TO-delivers m before m' .

Definition of Uniform

- A multicast is uniform if:

If any process delivers the multicast, all group members that don't fail will deliver it (even if the initial recipient fails immediately after delivery).


- Otherwise we say that the multicast is "not uniform"

Group Communication

Multicast Communication


- ❑ Group Communication is challenging even when all members of the group are static and aware of each other. **Dynamic groups**, with processes joining and leaving the group, are even more challenging. Most of the challenges are concerned with efficiency and delivery guarantees.
- ❑ **Efficiency** concerns include minimizing overhead activities and increasing throughput and bandwidth utilization.
- ❑ **Delivery guarantees** ensure that operations are completed.

Open and closed groups



Closed and Open Groups

- ❑ For multicast communications, a group is said to be closed if only members of the group can multicast to it. A process in a closed group sends to itself any messages to the group.
- ❑ A group is open if processes outside the group can send to it. Some algorithms assume closed groups while others assume open groups.



Summary

- ❑ Multicast communication can specify requirements for reliability and ordering, in terms of integrity, validity and agreement
- ❑ B-multicast
 - m a correct process will eventually deliver a message provided the multicaster does not crash
- ❑ reliable multicast
 - m in which the correct processes agree on the set of messages to be delivered;
 - m we showed two implementations: over B-multicast and IP multicast
- ❑ delivery ordering
 - m FIFO, total and causal delivery ordering.
 - m FIFO ordering by means of senders' sequence numbers
 - m total ordering by means of a sequencer or by agreement of sequence numbers between processes in a group
 - m causal ordering by means of vector timestamps
- ❑ the hold-back queue is a useful component in implementing multicast protocols