

Lecture 33


Map-Reduce

Map-Reduce

- ❑ Google processes 20 Petabytes/day
- ❑ Ebay 150 billion records/day
- ❑ Facebook 15 Terabytes/data
- ❑ 10,000 server cluster - 10 failures/day
- ❑ Era of "Big Data",

map (Functional Programming)

Creates a new list by applying f to each element of the input list; returns output in order.



```
fun map f []          = []  
  | map f (x::xs)    = (f x) :: (map f xs)
```

MapReduce Programming Model

- ❑ Programmers think in a *data-centric* fashion (transform the data)
- ❑ Have the framework handle the Hard Stuff:
 - Fault tolerance
 - Distributed execution, scheduling, concurrency
 - Coordination
 - Network communication

Basis API for MapReduce

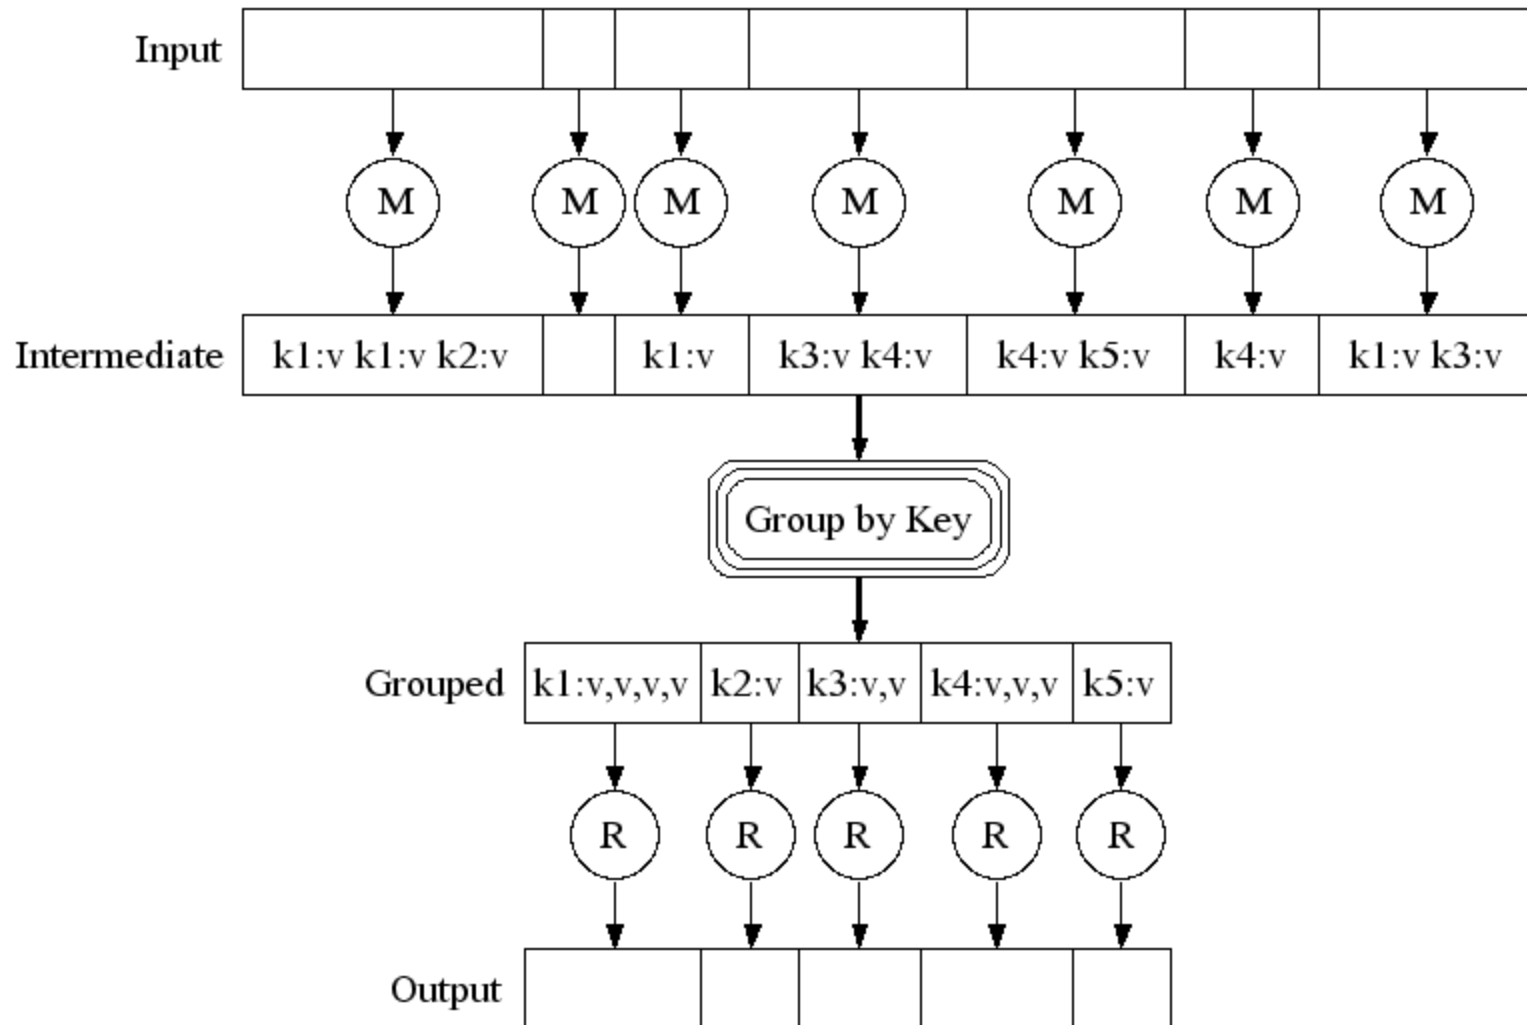
Map: $(\text{key1}, \text{value1}) \rightarrow [(\text{key2}, \text{value2})]$

Mapper is applied to every key-value pair that is input and computes one or more pairs of a new key with a new value. (for each key-value emits a list of pairs of key-values)

Reduce: $(\text{key2}, [\text{value2}]) \rightarrow [(\text{key3}, \text{value3})]$

For each key2 a reducer receives the key2 in sorted order and produces one or more pairs of a new key with a new value. (for each key2 emits a list of pairs of key-values)

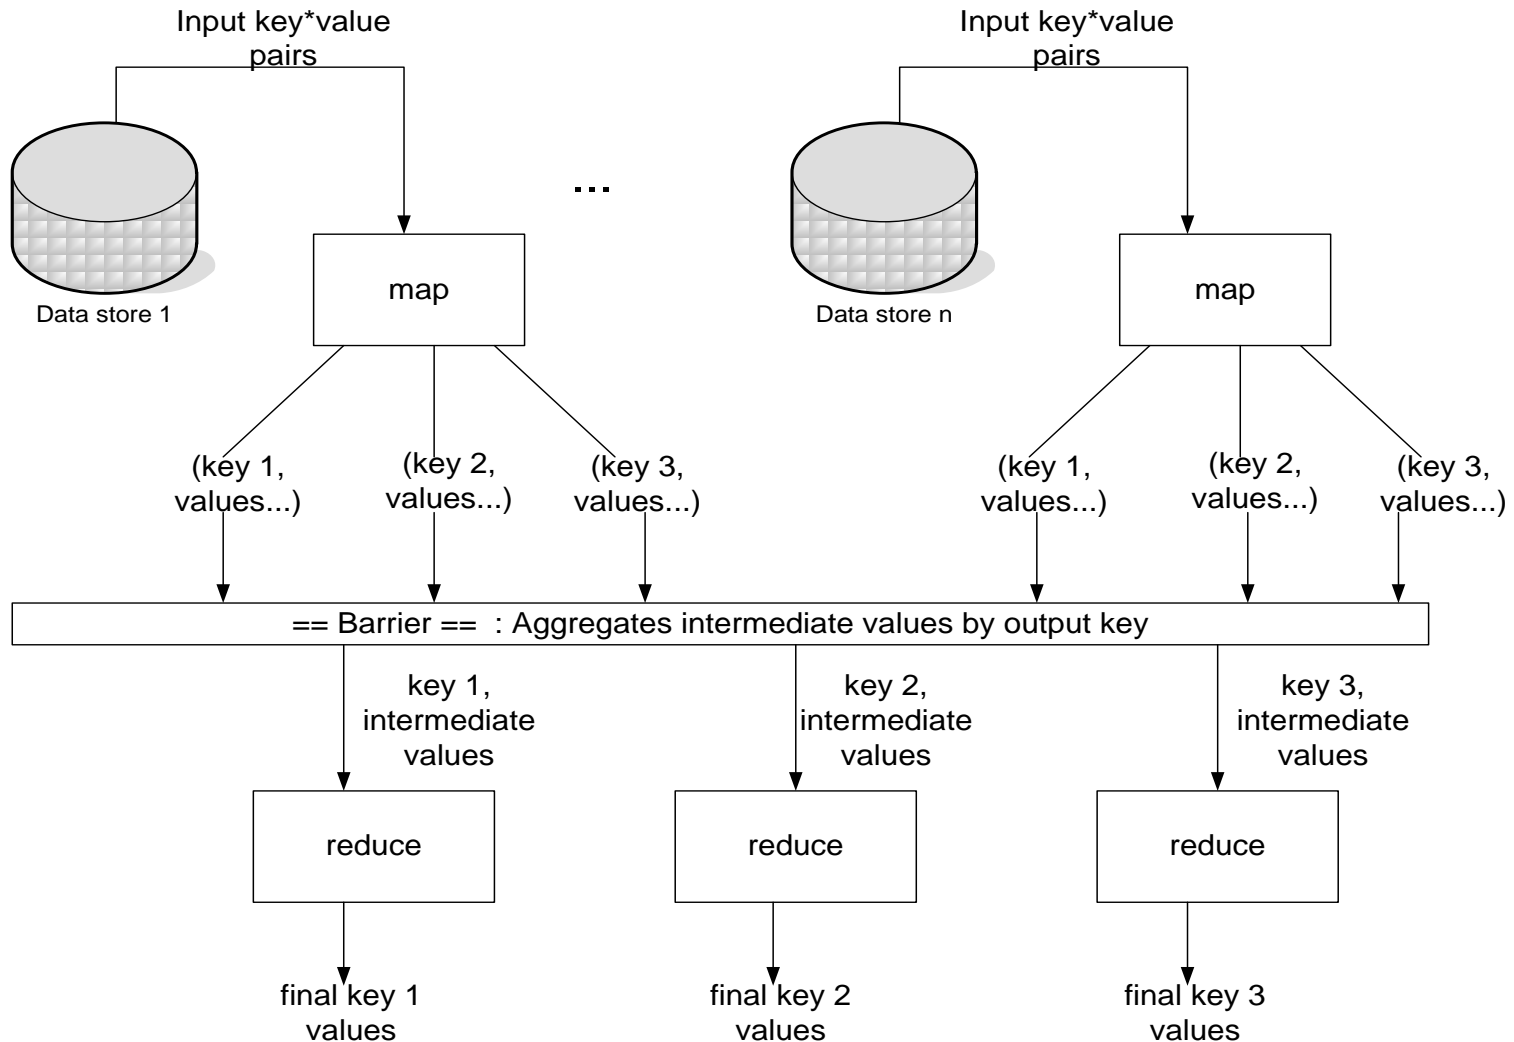
Map-Reduce



MapReduce System Model

- ❑ Designed for batch-oriented computations over large data sets (external memory algorithms)
 - ❑ Each map-reduce runs to completion before producing any output
 - ❑ Map-reduce output is written to stable storage
 - ❑ Map output to local disk, reduce output to HDFS
 - ❑ Designed to optimize “disk utilization”
 - ❑ Keep as many spindles active at once
- ❑ Simple, fault tolerance model
 - ❑ Can abort and restart a mapper or reducer.
 - ❑ check-point restart simple in this model

Map-Reduce System



Word count using MapReduce

Word count for a collection of documents:

Docid 1

This is a cat

Cat sits on a roof

Docid 2

The roof is a tin roof

There is a tin can on the roof

Docid 3

Cat kicks the can

It rolls on the roof and falls on the next roof

Docid 4

The cat rolls too

It sits on the can

Word Count using MapReduce

```
map(key, value):
```

```
    // key: document name; value: text of document  
    for each word w in value:  
        emit(w, 1)
```

```
reduce(key, values):
```

```
    // key: a word; values: iterator over counts  
    result = 0  
    for each count v in values:  
        result += v  
    emit(key, result)
```

Word Count using MapReduce:

Mapper

This is a cat

Cat sits on a roof

```
<this 1> <is 1> <a <1,1,>> <cat <1,1>> <sits 1> <on 1> <roof 1>
```

The roof is a tin roof

There is a tin can on the roof

```
<the <1,1>> <roof <1,1,1>> <is <1,1>> <a <1,1>> <tin <1,1>> <then 1>  
<can 1> <on 1>
```

Cat kicks the can

It rolls on the roof and falls on the next roof

```
<cat 1> <kicks 1> <the <1,1>> <can 1> <it 1> <roll 1> <on <1,1>>  
<roof <1,1>> <and 1> <falls 1> <next 1>
```

The cat rolls too

It sits on the can

```
<the <1,1>> <cat 1> <rolls 1> <too 1> <it 1> <sits 1> <on 1> <cat 1>
```

Combiners

- ❑ Often a map task will produce many pairs of the form $(k, v1), (k, v2), \dots$ for the same key k
 - E.g., popular words in Word Count
- ❑ Can save network time by pre-aggregating at mapper
 - $\text{combine}(k1, \text{list}(v1)) \rightarrow v2$
 - Usually same as reduce function
- ❑ Works only if reduce function is commutative and associative

Word Count using MapReduce: Combiner, Reducer

```
<this 1> <is 1> <a <1,1,>> <cat <1,1>> <sits 1> <on 1> <roof 1>  
<the <1,1>> <roof <1,1,1>> <is <1,1>> <a <1,1>> <tin <1,1>> <then 1>  
  <can 1> <on 1>  
<cat 1> <kicks 1> <the <1,1>> <can 1> <it 1> <roll 1> <on <1,1>>  
  <roof <1,1>> <and 1> <falls 1> <next 1>  
<the <1,1>> <cat 1> <rolls 1> <too 1> <it 1> <sits 1> <on 1> <cat 1>
```

Combine the counts of all the same words:

```
<cat <1,1,1,1>>  
<roof <1,1,1,1,1,1>>  
<can <1, 1,1>>
```

...

Reduce (sum in this case) the counts:

```
<cat 4>  
<can 3>  
<roof 6>
```

Map-Reduce

- ❑ Mapper
- ❑ Reducers
- ❑ Combiners

Data-intensive text processing with MapReduce

Jimmy Lin and Chris Dyer.

Book, Online in English, 2010

(<http://lintool.github.com/MapReduceAlgorithms/>)

MapReduce/Hadoop

- ❑ Around 2008, Yahoo! developed the open-source variant of MapReduce named Hadoop
- ❑ After 2008, MapReduce/Hadoop become a key technology component in cloud computing
- ❑ In 2010, the U.S. conferred the MapReduce patent to Google



MapReduce

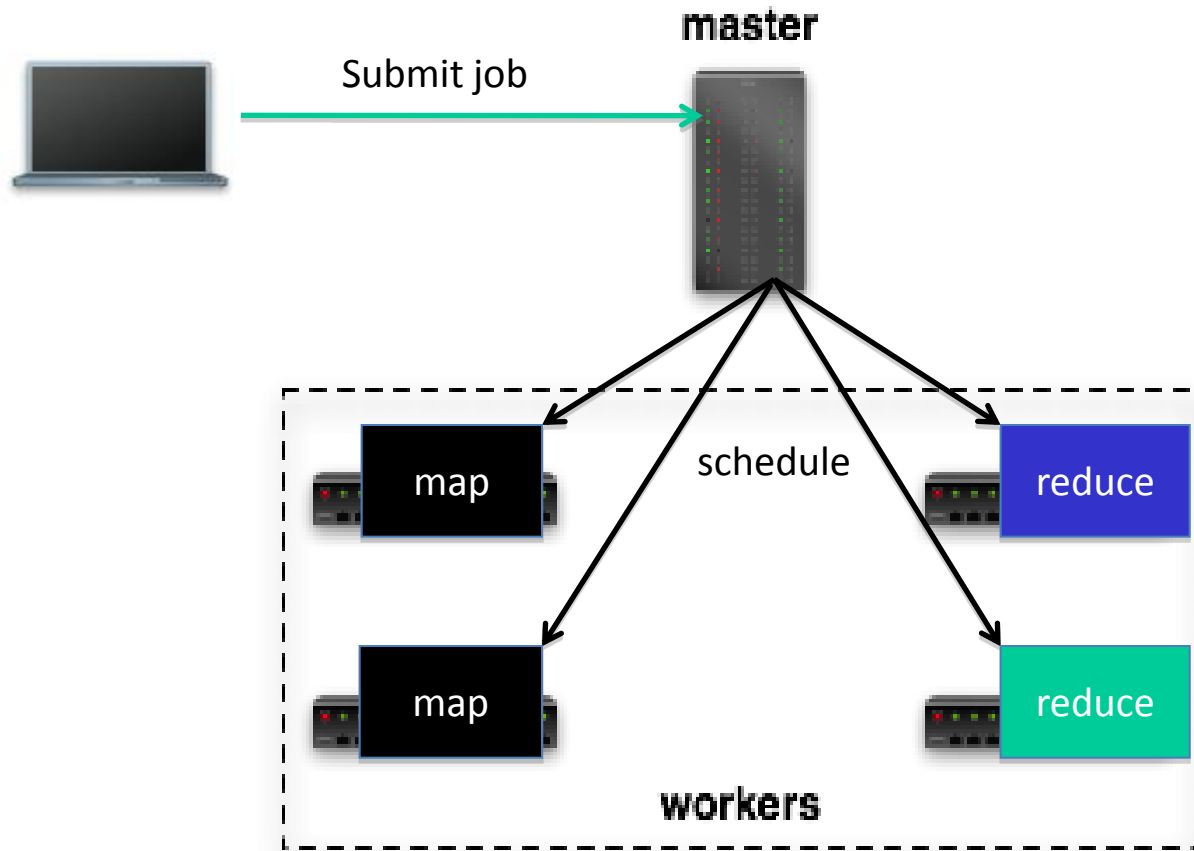


Hadoop




... Hadoop or variants ...


Dataflow in Hadoop




Dataflow in Hadoop



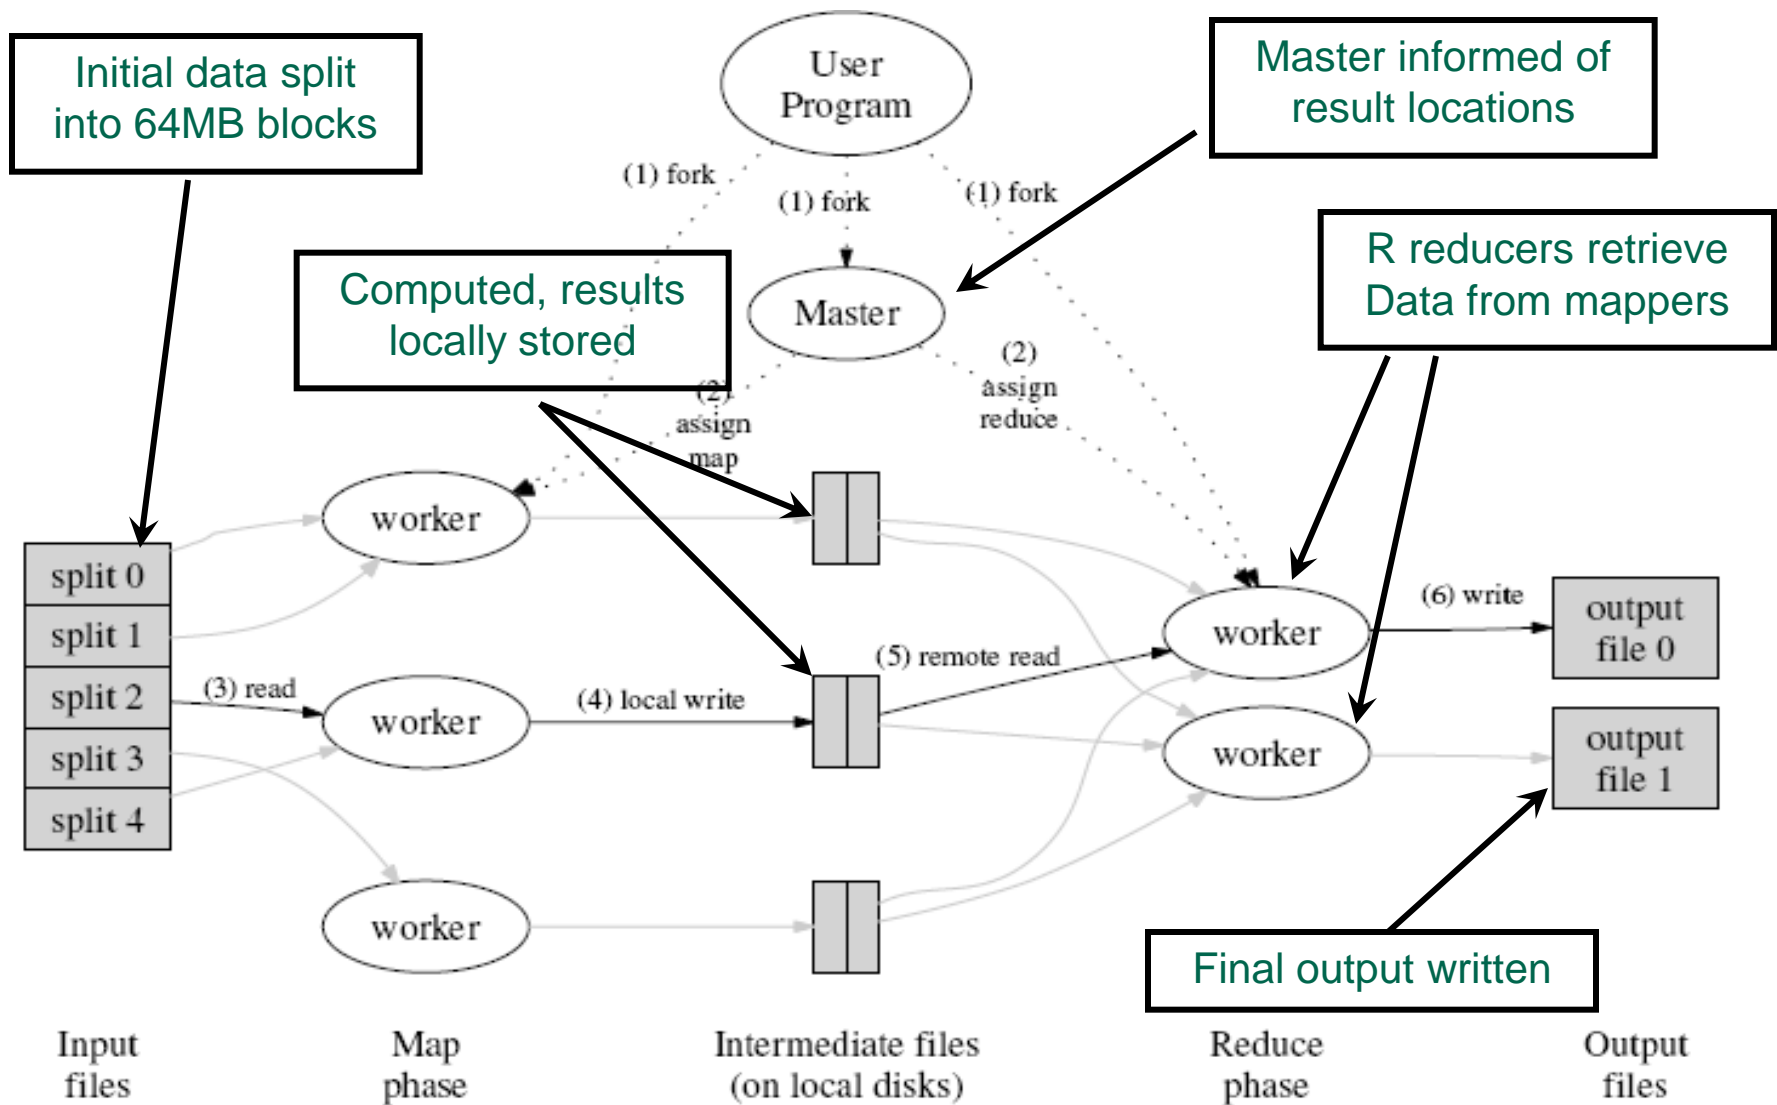
Dataflow in Hadoop



Dataflow in Hadoop



Workflow



Critique

MapReduce: A major step backwards

-- David J. DeWitt and Michael Stonebraker

(MapReduce) is

- A giant step backward in the programming paradigm for large-scale data intensive applications
- A sub-optimal implementation, in that it uses brute force instead of indexing
- Not novel at all
- Missing features
- Incompatible with all of the tools DBMS users have come to depend on

Questions

- ❑ MapReduce provides an easy-to-use framework for parallel programming, but is it the most efficient and best solution to program execution in datacenters?
- ❑ Comparison to database systems:
 - MapReduce is not a database system, nor a query language
 - It is possible to use MapReduce to implement some of the parallel query processing functions
 - What are the real limitations?
 - Scalability

Limitations

- ❑ Inefficient for general programming (and not designed for that)
 - Hard to handle data with complex dependence, frequent updates, etc.
 - High overhead, bursty I/O, streaming data
- ❑ Difficult to optimize (limited opportunities)
- ❑ Hadoop
 - Typically orders of magnitude slower than a parallel computing solution
 - BUT, environment is different because it depends on file-system (designed for scalability!)

Limitations

- ❑ Proprietary solution developed in an environment with one prevailing application (web search)
 - The assumptions introduce several important constraints in data and logic
 - Not a general-purpose parallel execution technology
- ❑ Design choices in MapReduce
 - Optimizes for throughput rather than latency
 - Optimizes for large data set rather than small data structures (external data algorithms)
 - Optimizes for coarse-grained parallelism rather than fine-grained

Presentations

Josh Carter, http://multipart-mixed.com/software/mapreduce_presentation.pdf

Ralf Lammel, *Google's MapReduce Programming Model - Revisited*
<http://code.google.com/edu/parallel/mapreduce-tutorial.htm>

References

- ❑ [Al-Fares] Al-Fares, M., Loukissas, A., and Vahdat, A. A scalable, commodity data center network architecture. In Proceedings of the ACM SIGCOMM 2008 Conference on Data Communication (Seattle, WA, USA, August 17 - 22, 2008). SIGCOMM '08. 63-74.
- ❑ [Andersen] David G. Andersen, Jason Franklin, Michael Kaminsky, Amar Phanishayee, Lawrence Tan, Vijay Vasudevan. FAWN: A Fast Array of Wimpy Nodes. SOSP'09.
- ❑ [Barraso] Luiz Barroso, Jeffrey Dean, Urs Hoelzle, "Web Search for a Planet: The Google Cluster Architecture," IEEE Micro, vol. 23, no. 2, pp. 22-28, Mar./Apr. 2003
- ❑ [Ghemawat] Ghemawat, S., Gobioff, H., and Leung, S. 2003. The Google file system. In Proceedings of the Nineteenth ACM Symposium on Operating Systems Principles (Bolton Landing, NY, USA, October 19 - 22, 2003). SOSP '03. ACM, New York, NY, 29-43.
- ❑ [Guo 2009] Chuanxiong Guo, Guohan Lu, Dan Li, Xuan Zhang, Haitao Wu, Yunfeng Shi, Chen Tian, Yongguang Zhang, and Songwu Lu, BCube: A High Performance, Server-centric Network Architecture for Modular Data Centers, in ACM SIGCOMM 09.

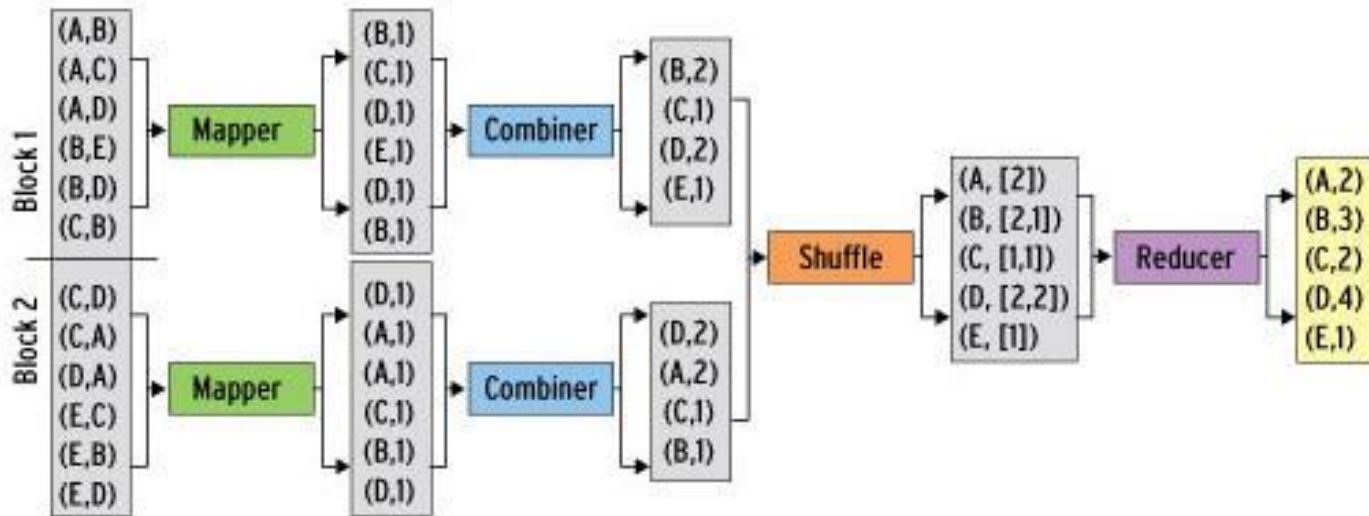


- ❑ [Chang] Chang, F., Dean, J., Ghemawat, S., Hsieh, W. C., Wallach, D. A., Burrows, M., Chandra, T., Fikes, A., and Gruber, R. E. Bigtable: a distributed storage system for structured data. In Proceedings of the 7th Symposium on Operating Systems Design and Implementation (Seattle, Washington, November 06 - 08, 2006). 205-218.
- ❑ [Cooper] Cooper, B. F., Ramakrishnan, R., Srivastava, U., Silberstein, A., Bohannon, P., Jacobsen, H., Puz, N., Weaver, D., and Yerneni, R. PNUTS: Yahoo!'s hosted data serving platform. Proc. VLDB Endow. 1, 2 (Aug. 2008), 1277-1288.
- ❑ [Dean] Dean, J. and Ghemawat, S. 2004. MapReduce: simplified data processing on large clusters. In Proceedings of the 6th Conference on Symposium on Operating Systems Design & Implementation - Volume 6 (San Francisco, CA, December 06 - 08, 2004).
- ❑ [Ghemawat] Ghemawat, S., Gobioff, H., and Leung, S. 2003. The Google file system. In Proceedings of the Nineteenth ACM Symposium on Operating Systems Principles (Bolton Landing, NY, USA, October 19 - 22, 2003). SOSP '03. ACM, New York, NY, 29-43.
- ❑ [Guo 2009] Chuanxiong Guo, Guohan Lu, Dan Li, Xuan Zhang, Haitao Wu, Yunfeng Shi, Chen Tian, Yongguang Zhang, and Songwu Lu, BCube: A High Performance, Server-centric Network Architecture for Modular Data Centers, in ACM SIGCOMM 09.

Combiners

- ❑ Often a map task will produce many pairs of the form $(k, v1), (k, v2), \dots$ for the same key k
 - E.g., popular words in Word Count
- ❑ Can save network time by pre-aggregating at mapper
 - $\text{combine}(k1, \text{list}(v1)) \rightarrow v2$
 - Usually same as reduce function
- ❑ Works only if reduce function is commutative and associative

Word Count Example

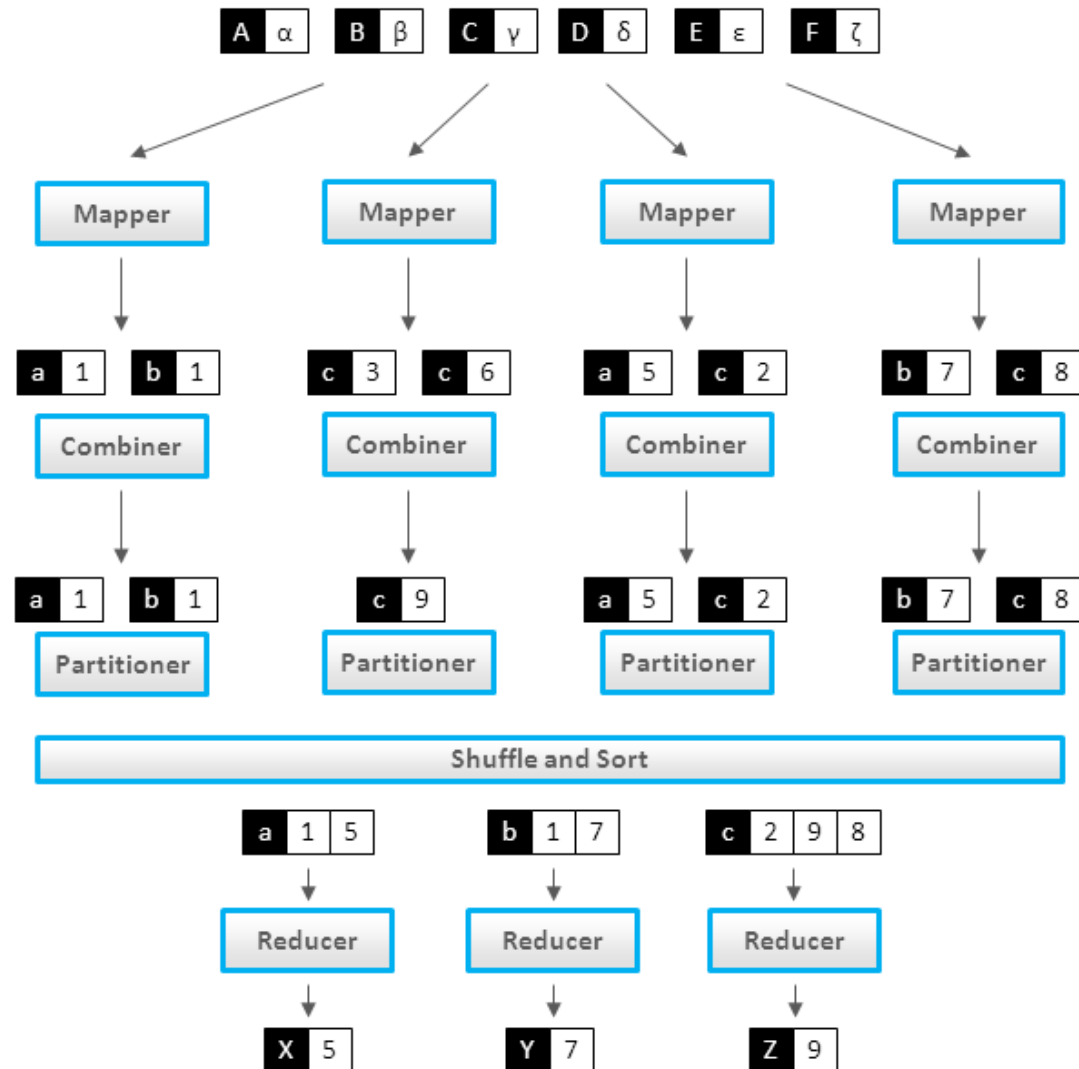


Partition Function

- ❑ Inputs to map tasks are created by contiguous splits of input file
- ❑ For reduce, we need to ensure that records with the same intermediate key end up at the same worker
- ❑ System uses a default partition function e.g., $\text{hash}(\text{key}) \bmod R$
- ❑ Sometimes useful to override
 - E.g., $\text{hash}(\text{hostname}(\text{URL})) \bmod R$ ensures URLs from a host end up in the same output file

Map-Reduce

Word Count using MapReduce



Text Processing --Co-occurrence

Co-occurrence of two words in a document, in general co-occurrence with respect to some **context**

Docid 1

This is a cat

Cat sits on a roof

Docid 2

The roof is a tin roof

There is a tin can on the roof

[illegible]

Relative Frequency

Frequency of one co-occurrence with all the co-occurrence:

$$f(w_j | w_i) = \frac{N(w_i, w_j)}{\sum_{w'} N(w_i, w')}$$

Relative Frequency using MapReduce

Collection of documents:

Docid 1

This is a cat

Cat sits on a roof

Docid 2

The roof is a tin roof

There is a tin can on the roof

Docid 3

Cat kicks the can

It rolls on the roof and falls on the next roof

Docid 4

The cat rolls too

It sits on the can

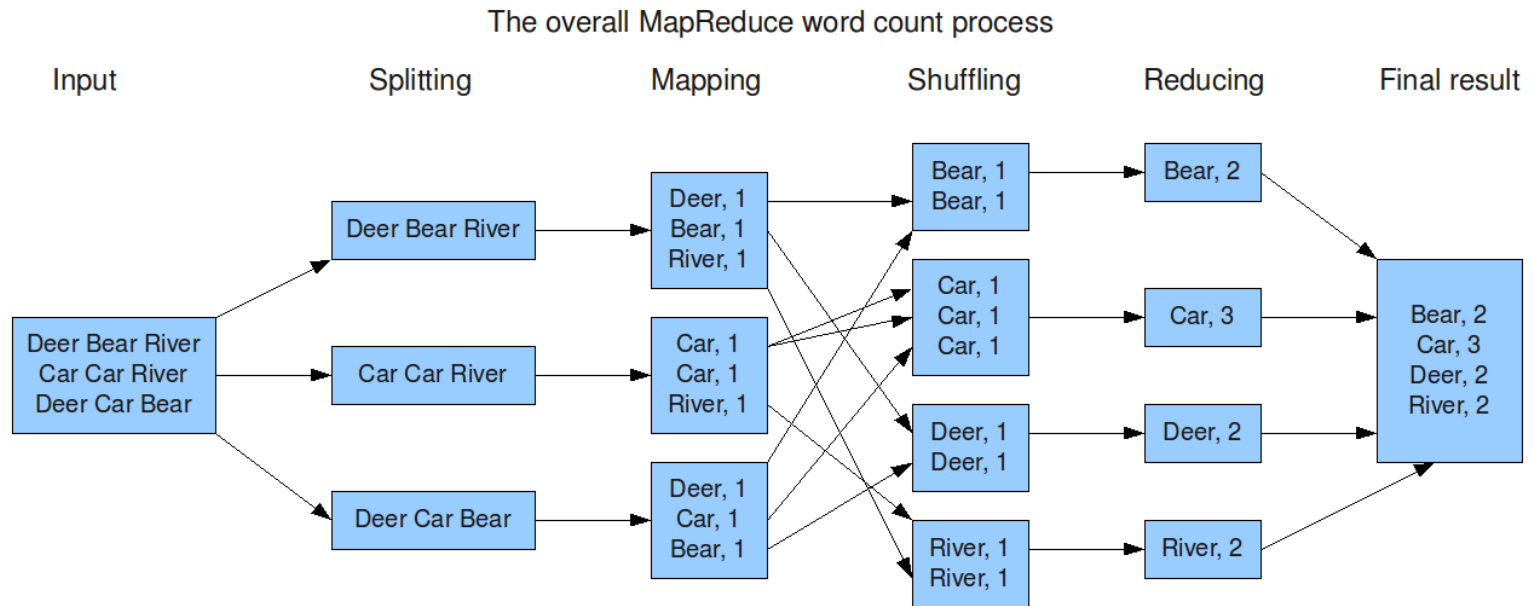
Order Inversion

- ❑ Emitting a special key-value for each pair to capture the contribution to the sum
- ❑ Controlling the sort order of the intermediate key so that the sum communication appears before the pairs
- ❑ Defining a custom partitioner to ensure the pairs are sorted in lexicographical order and appear at the same reducer
- ❑ Preserve state across multiple keys in the reducer to first compute total sum and then dividing this with the co-occurrence value.

Data-Warehousing (OLAP)

- ❑ One-to-one Joins
 - Reduce-side joining
- ❑ One-to-many Joins
- ❑ Many-to-many Joins
- ❑ Simple merging with the mapper

Word Count Example



Word Count using MapReduce

```
map(key, value):
```

```
    // key: document name; value: text of document  
    for each word w in value:  
        emit(w, 1)
```

```
reduce(key, values):
```

```
    // key: a word; values: iterator over counts  
    result = 0  
    for each count v in values:  
        result += v  
    emit(key, result)
```